CS 4530 Software Engineering

Module 14: Principles and Patterns of Cloud Infrastructure

Adeel Bhutta, Joydeep Mitra and Mitch Wand

Khoury College of Computer Sciences

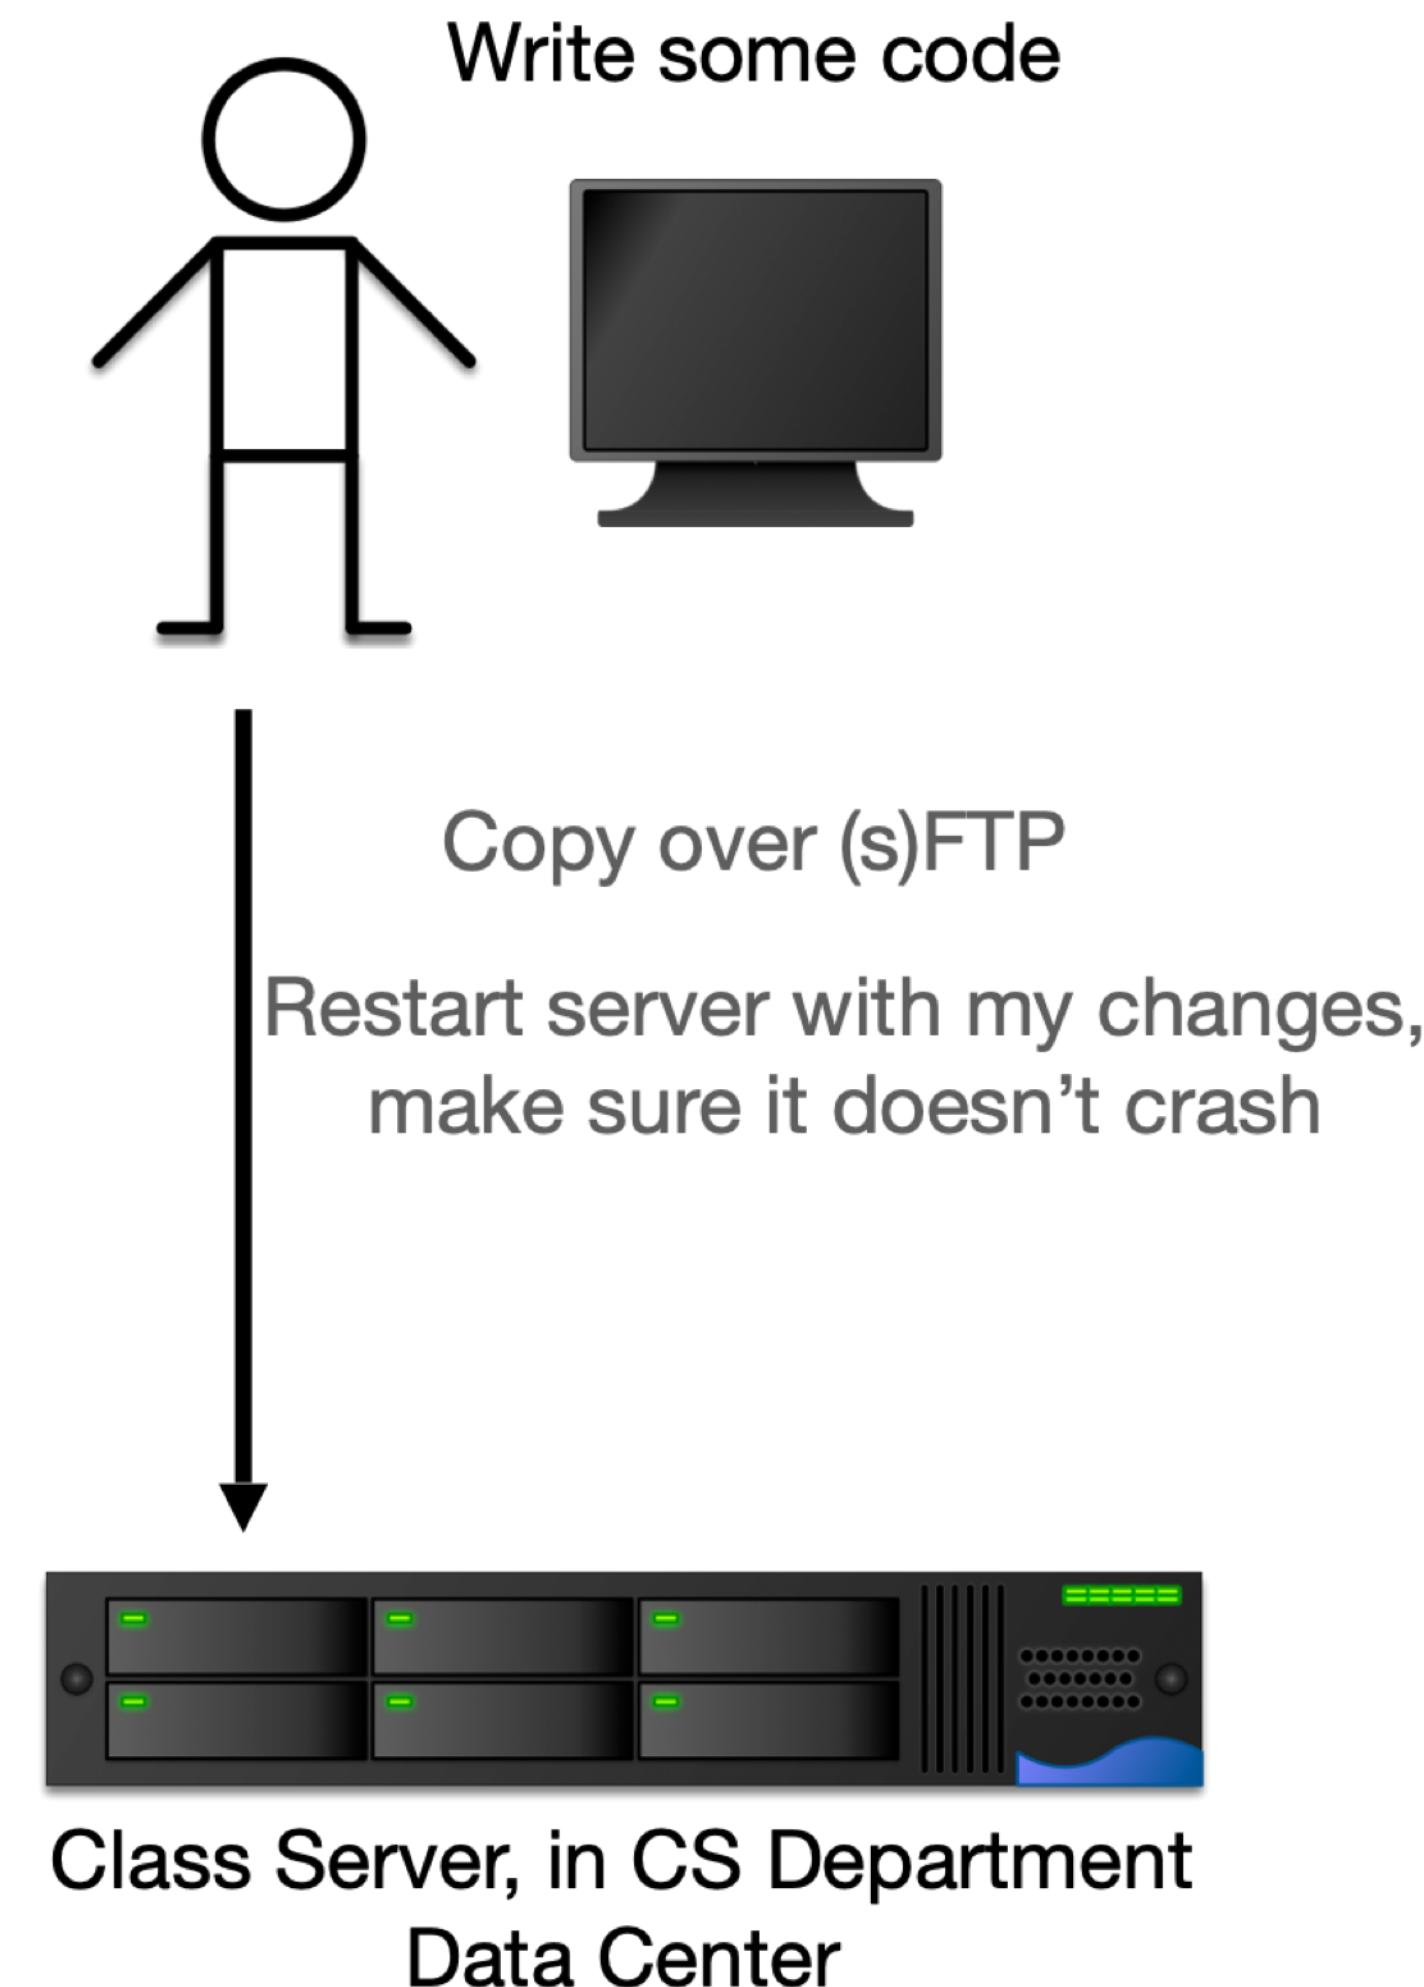# Learning objectives for this lesson

- By the end of this lesson, you should be able to…

  - Explain what "cloud" computing is and why it is important

  - Explain why shared infrastructure is important in cloud computing

  - Describe the difference between virtual machines and containers

  - Discuss trade-offs that you might consider for self or vendor-managed platforms
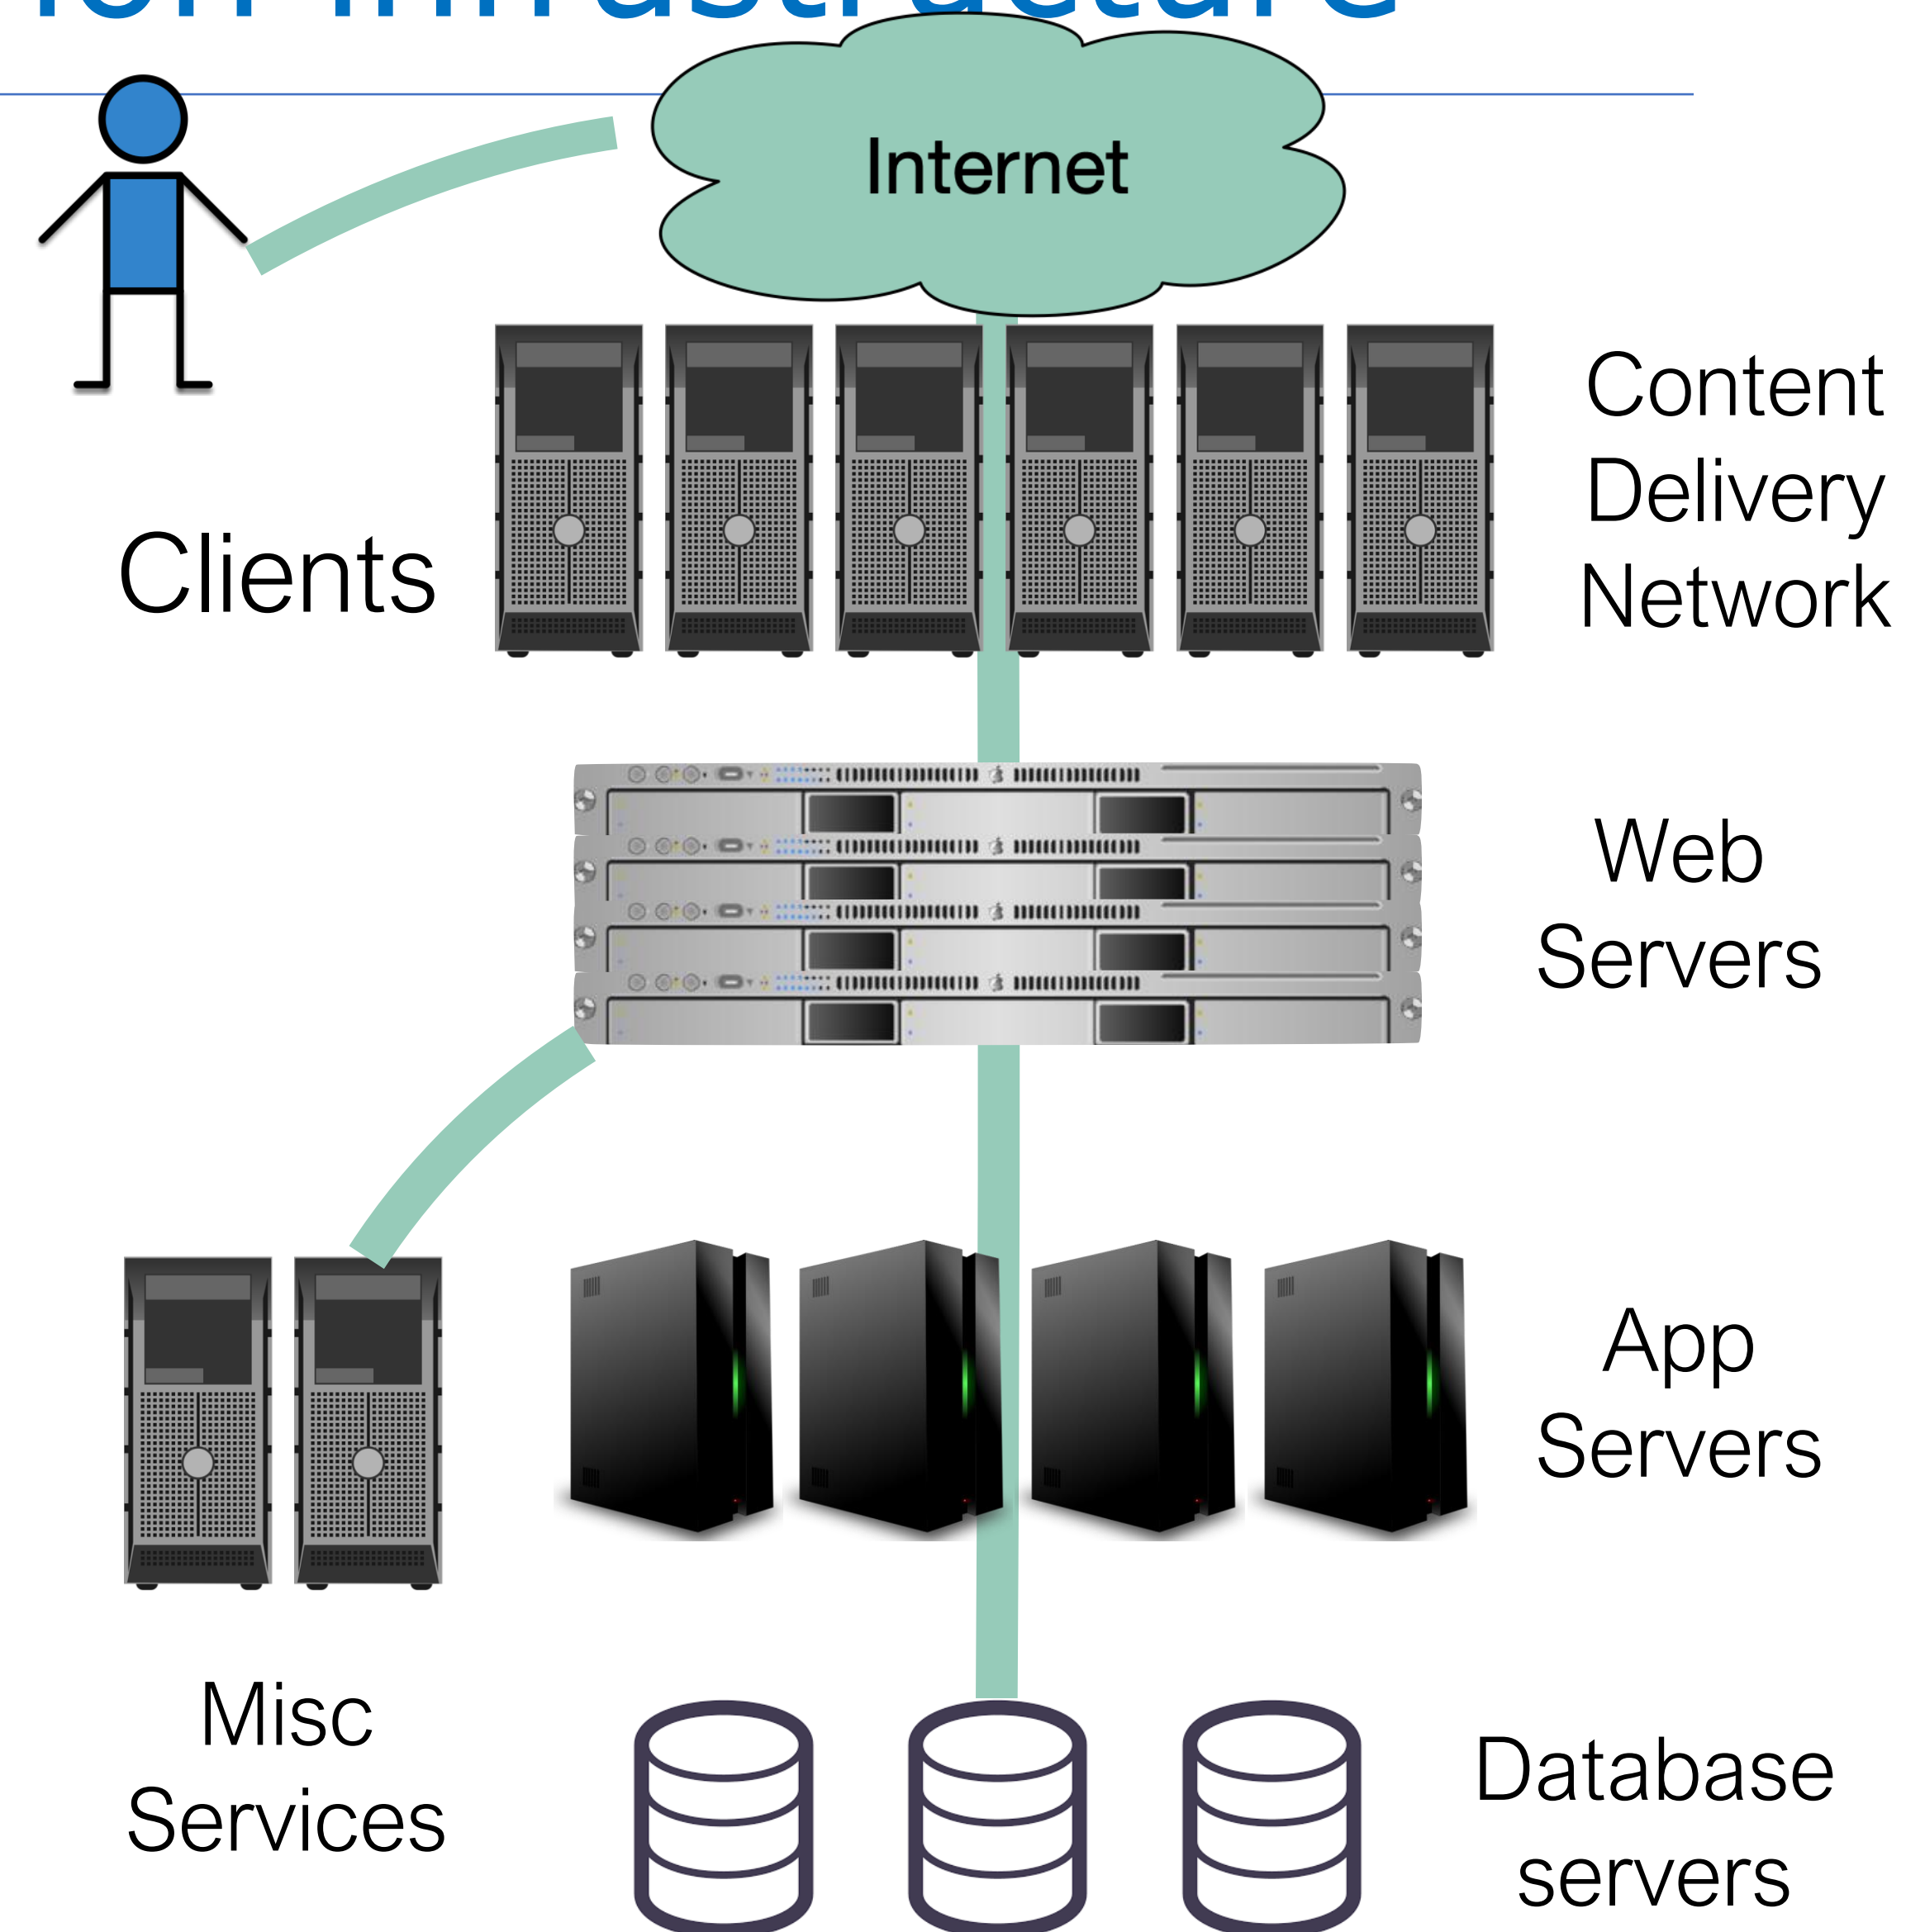
How to deploy web apps?

- What we need:
  - A server that can run our application
  - A network that is configured to route requests from an address to that server
- Questions to think about:
  - What software do we need to run besides our application code? (Databases, caches, etc?)
  - Where does this server come from? (Buy/Borrow?)
  - Who else gets to use this server? (Multi-tenancy or exclusive?)
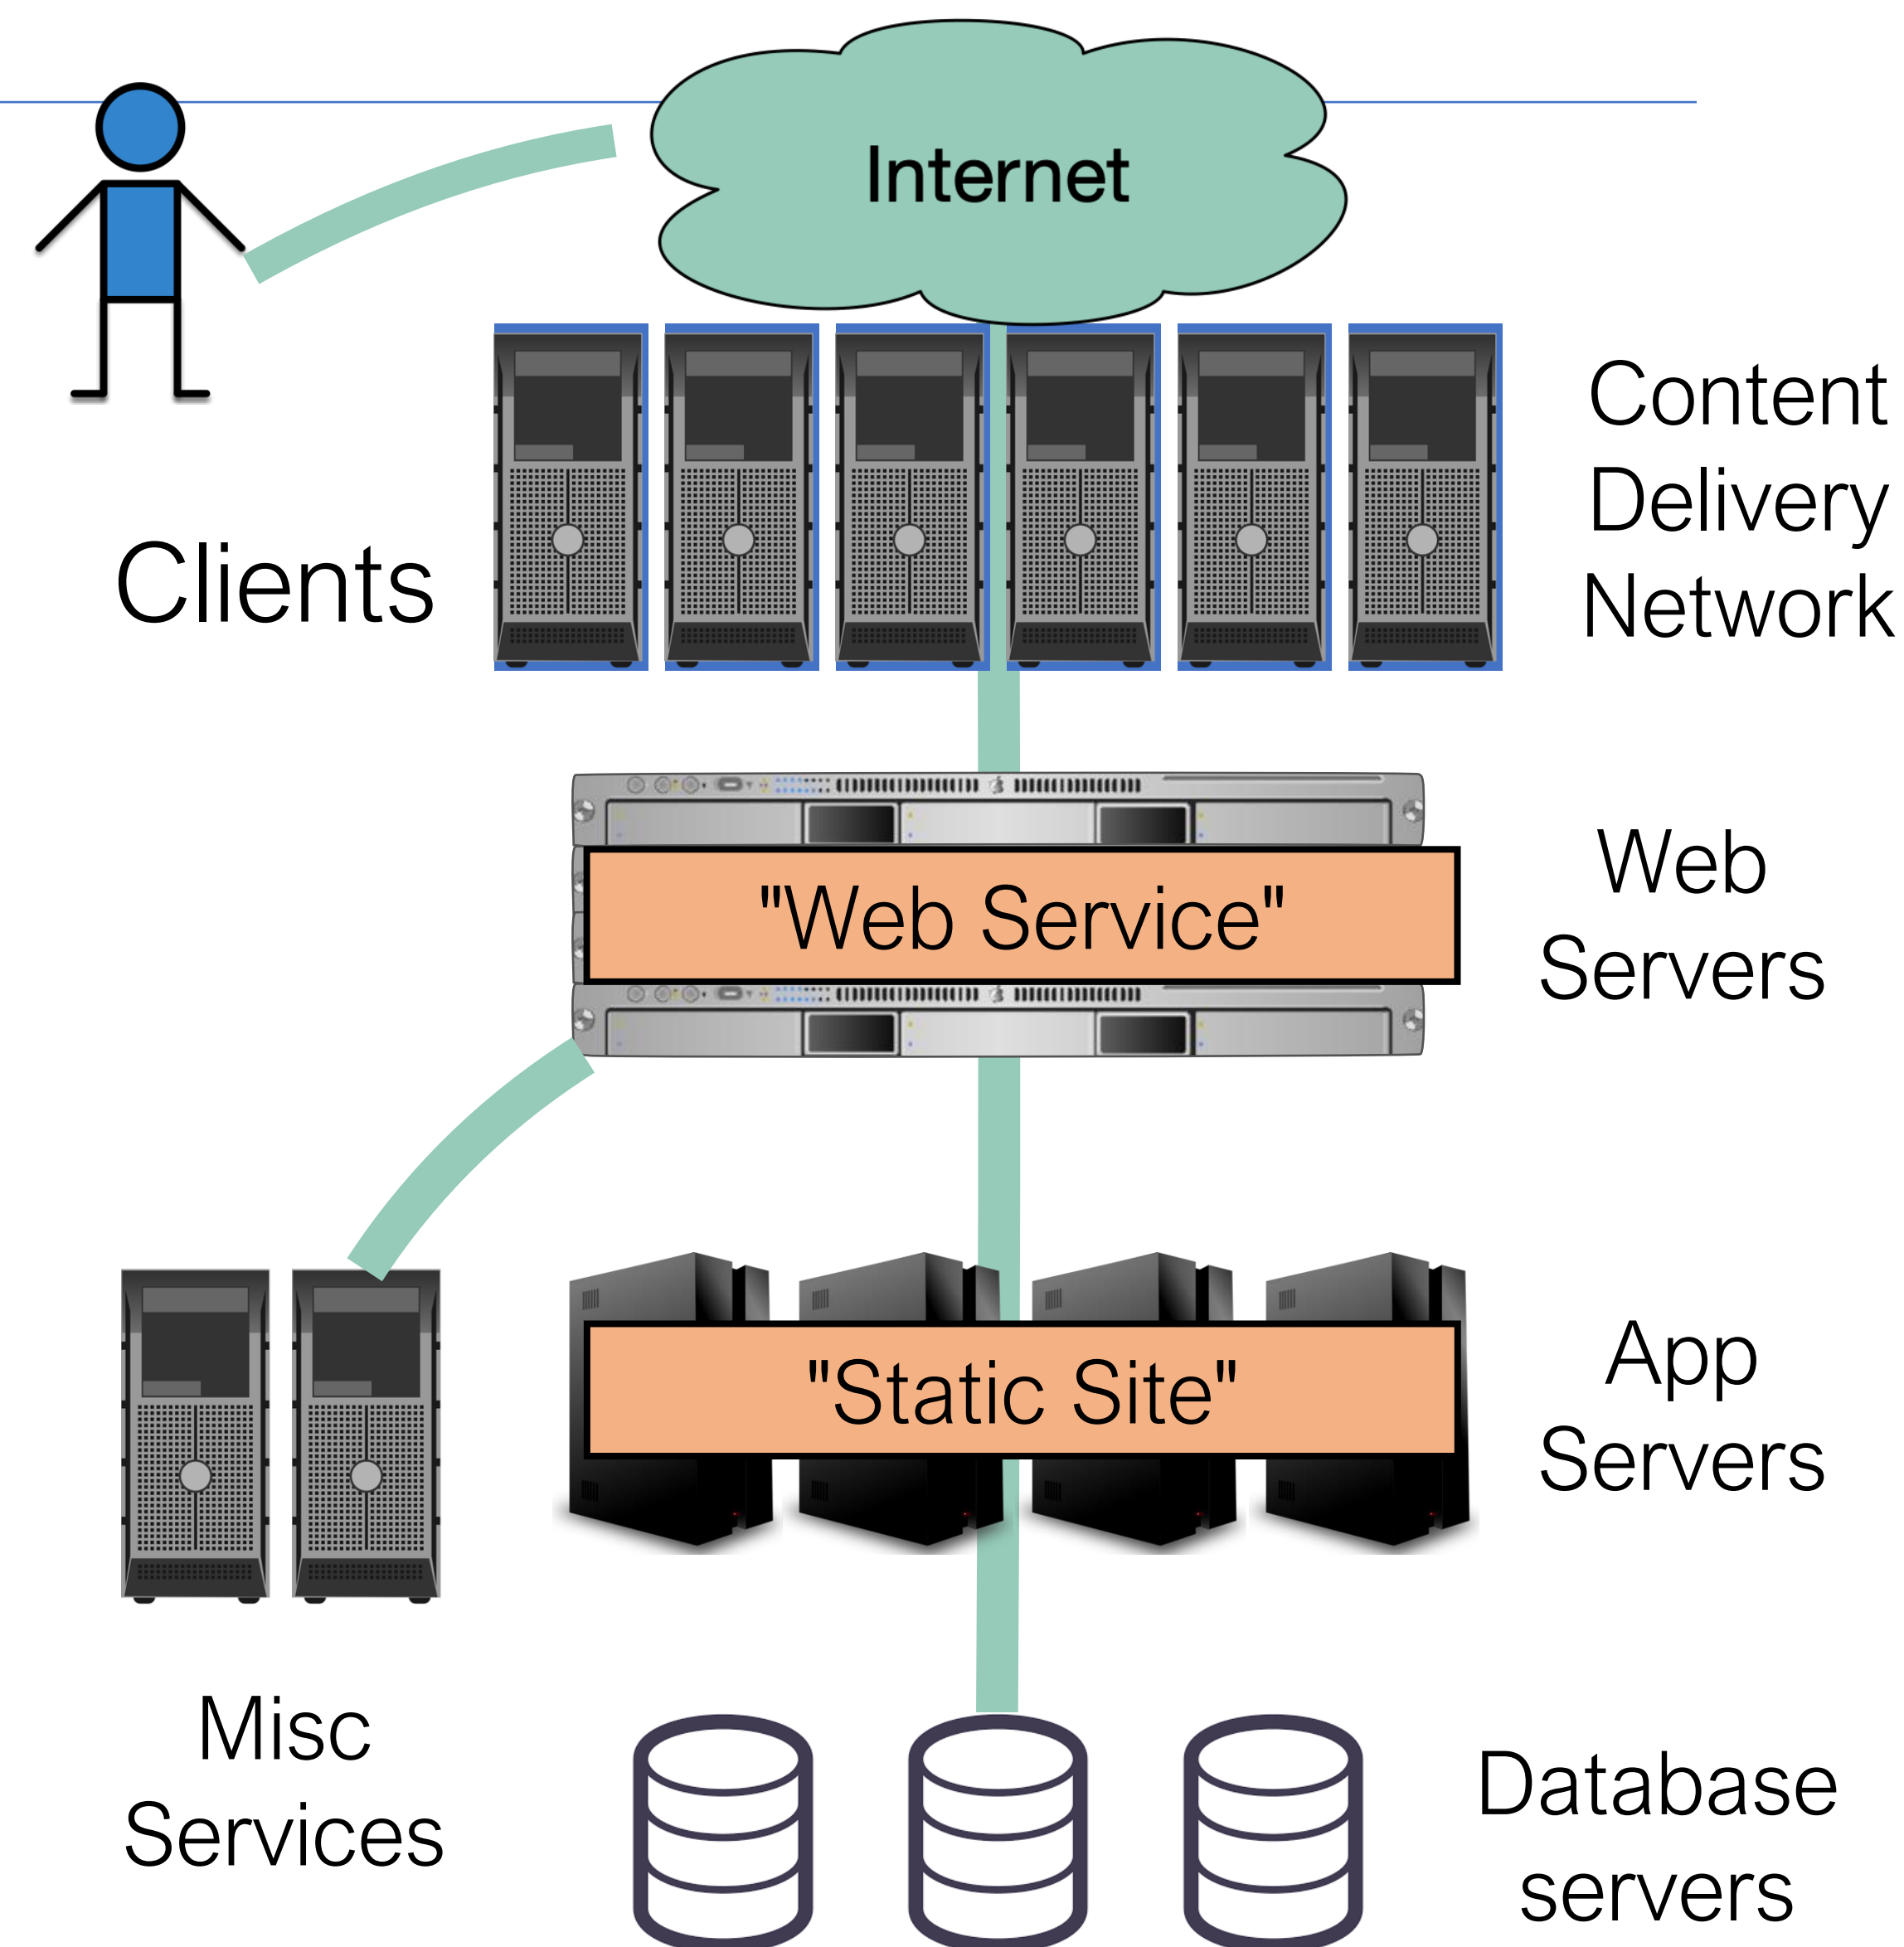 - Who maintains the server and software? (Updates OS, libraries, etc?)

Write some code

Copy over (s)FTP

Restart server with my changes, make sure it doesn't crash

Class Server, in CS Department Data Center

Many apps rely on common infrastructure

- Content delivery network: caches static content "at the edge" (e.g. cloudflare, Akamai)

- Web servers: Speak HTTP, serve static content (eg REACT), load balance between app servers (e.g. haproxy, traefik)

- App servers: Run our backend application (e.g. nodejs)

- Misc services: Logging, monitoring, firewall

- Database servers: Persistent data

Clients

Internet

Content Delivery Network

Web Servers

App Servers

Misc Services

Database servers

Render.com terminology

- Content delivery network: caches static content "at the edge" (e.g. cloudflare, Akamai)

- Web servers: Speak HTTP, serve static content (eg REACT), load balance between app servers (e.g. haproxy, traefik)

- App servers: Run our backend application (e.g. nodejs)

- Misc services: Logging, monitoring, firewall
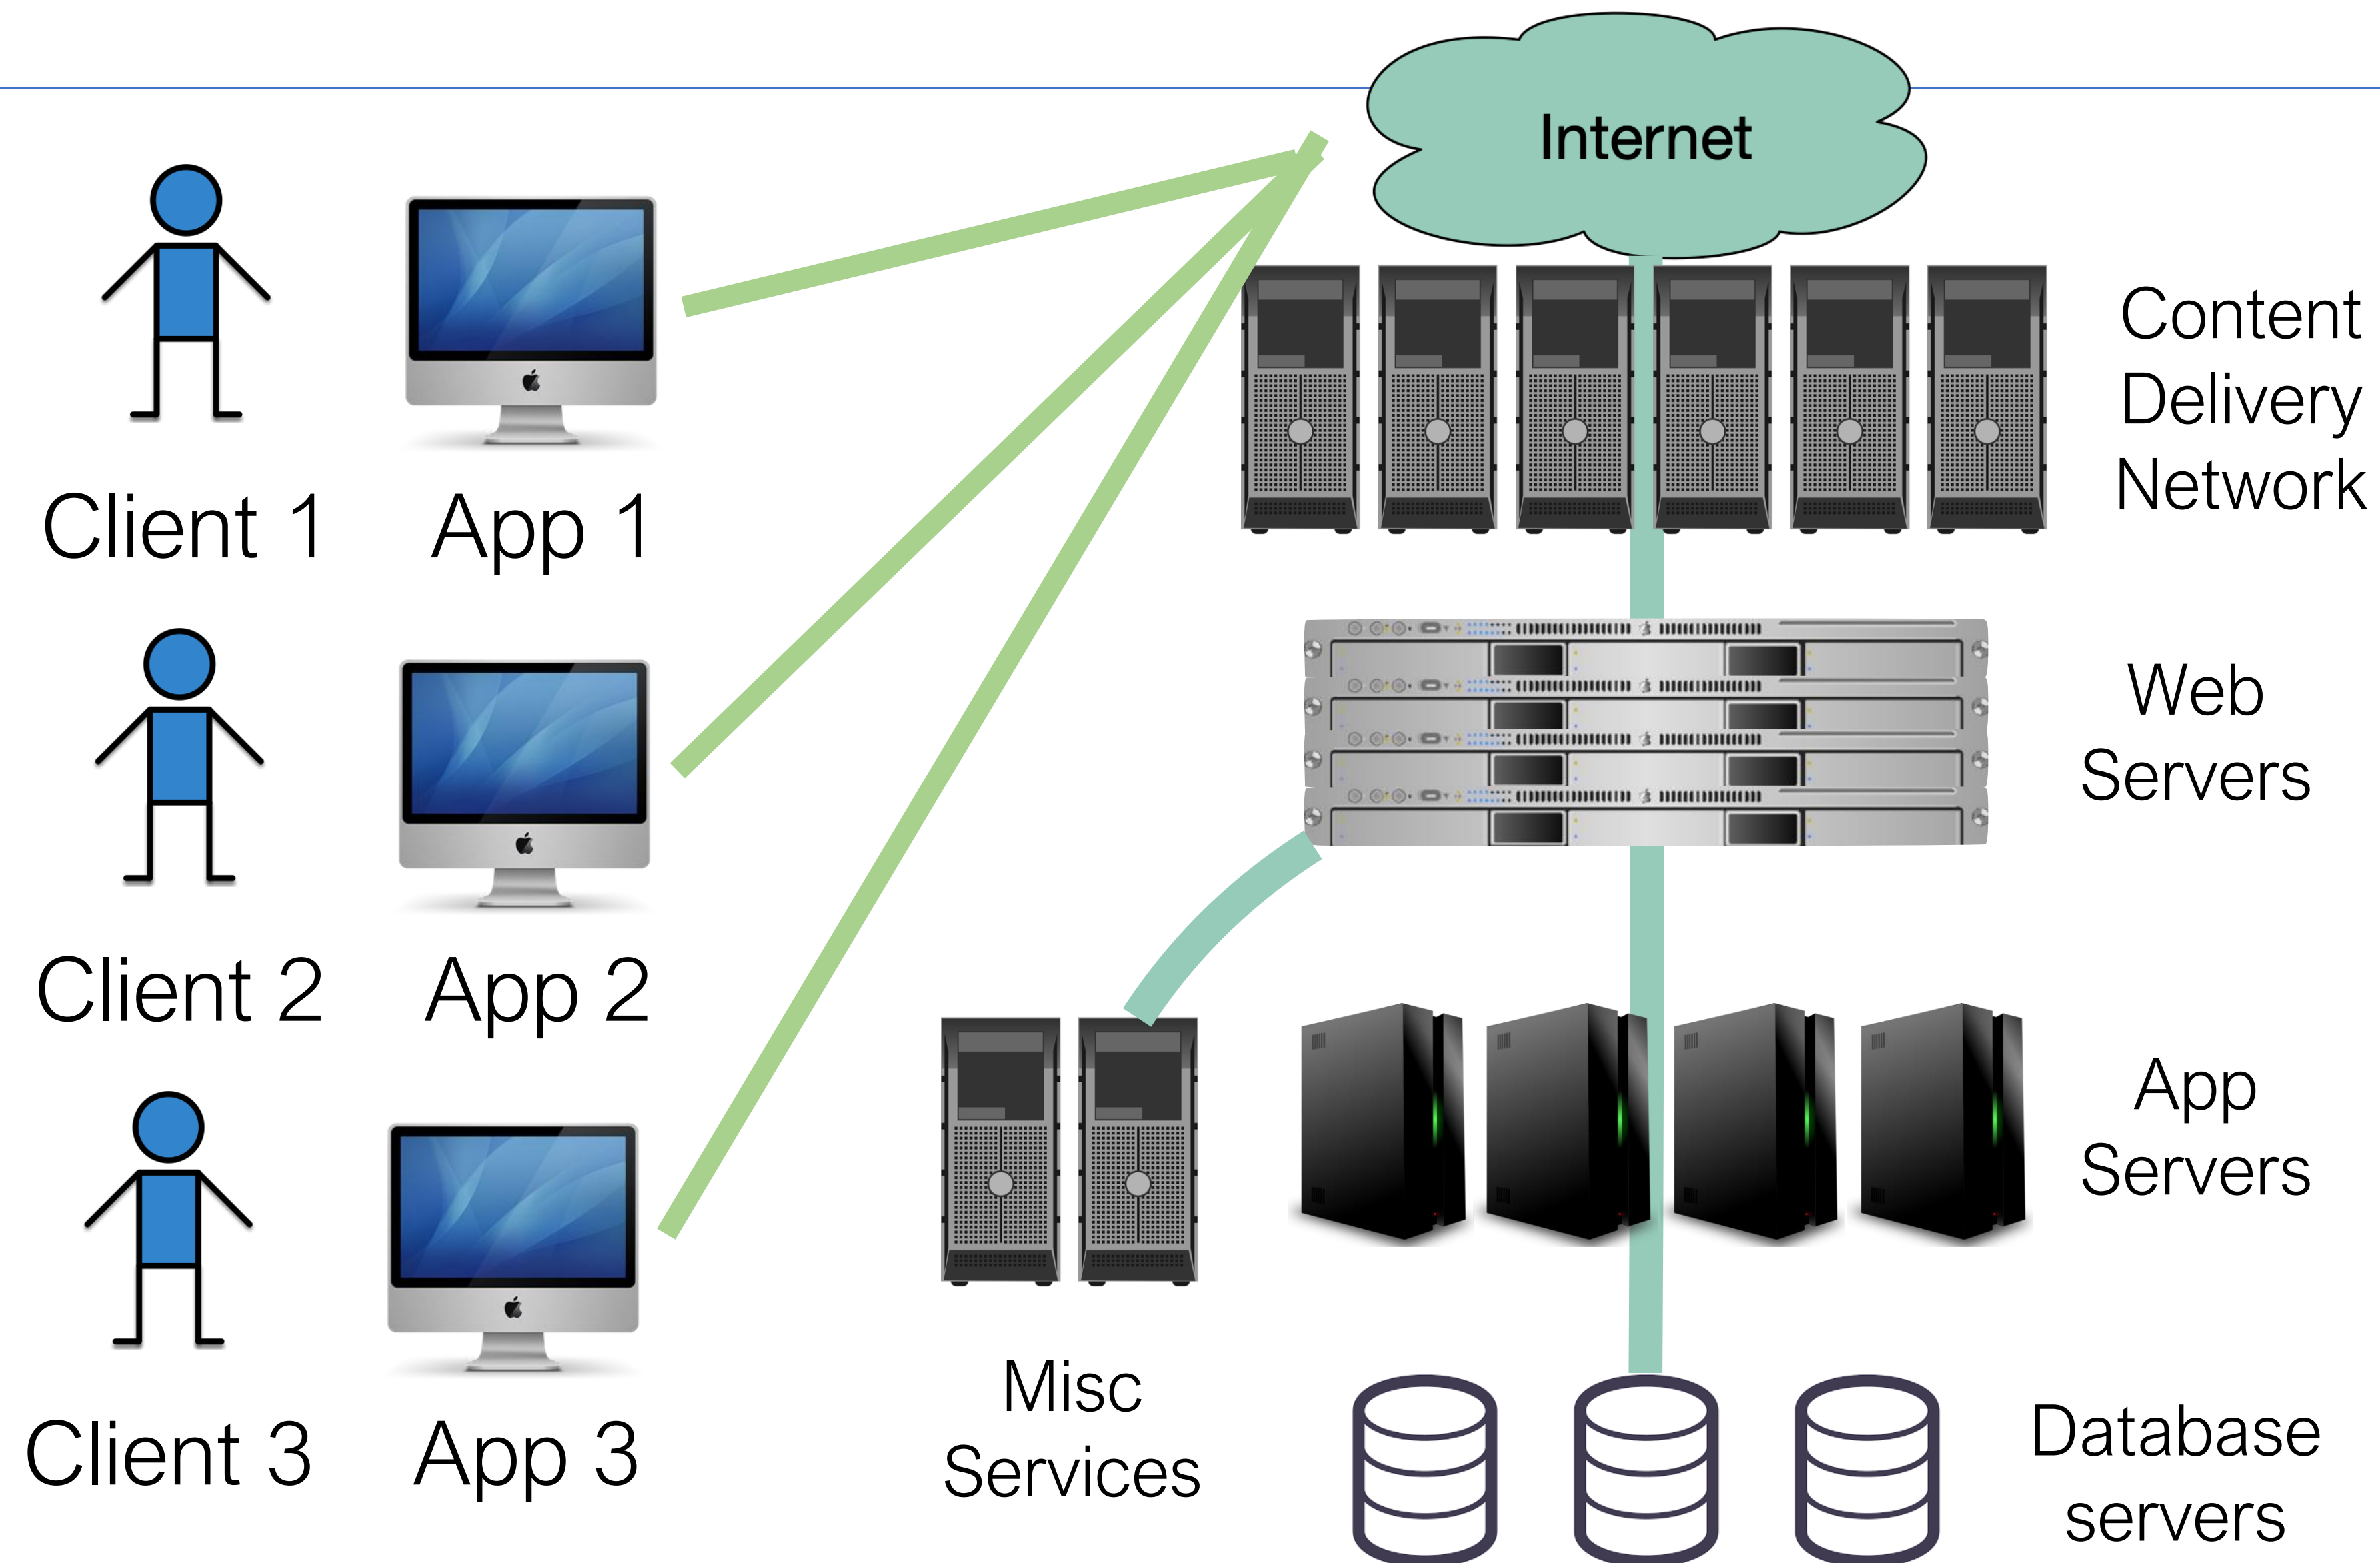
- Database servers: Persistent data

Clients

Internet

Content Delivery Network

"Web Service"

Web Servers

"Static Site"
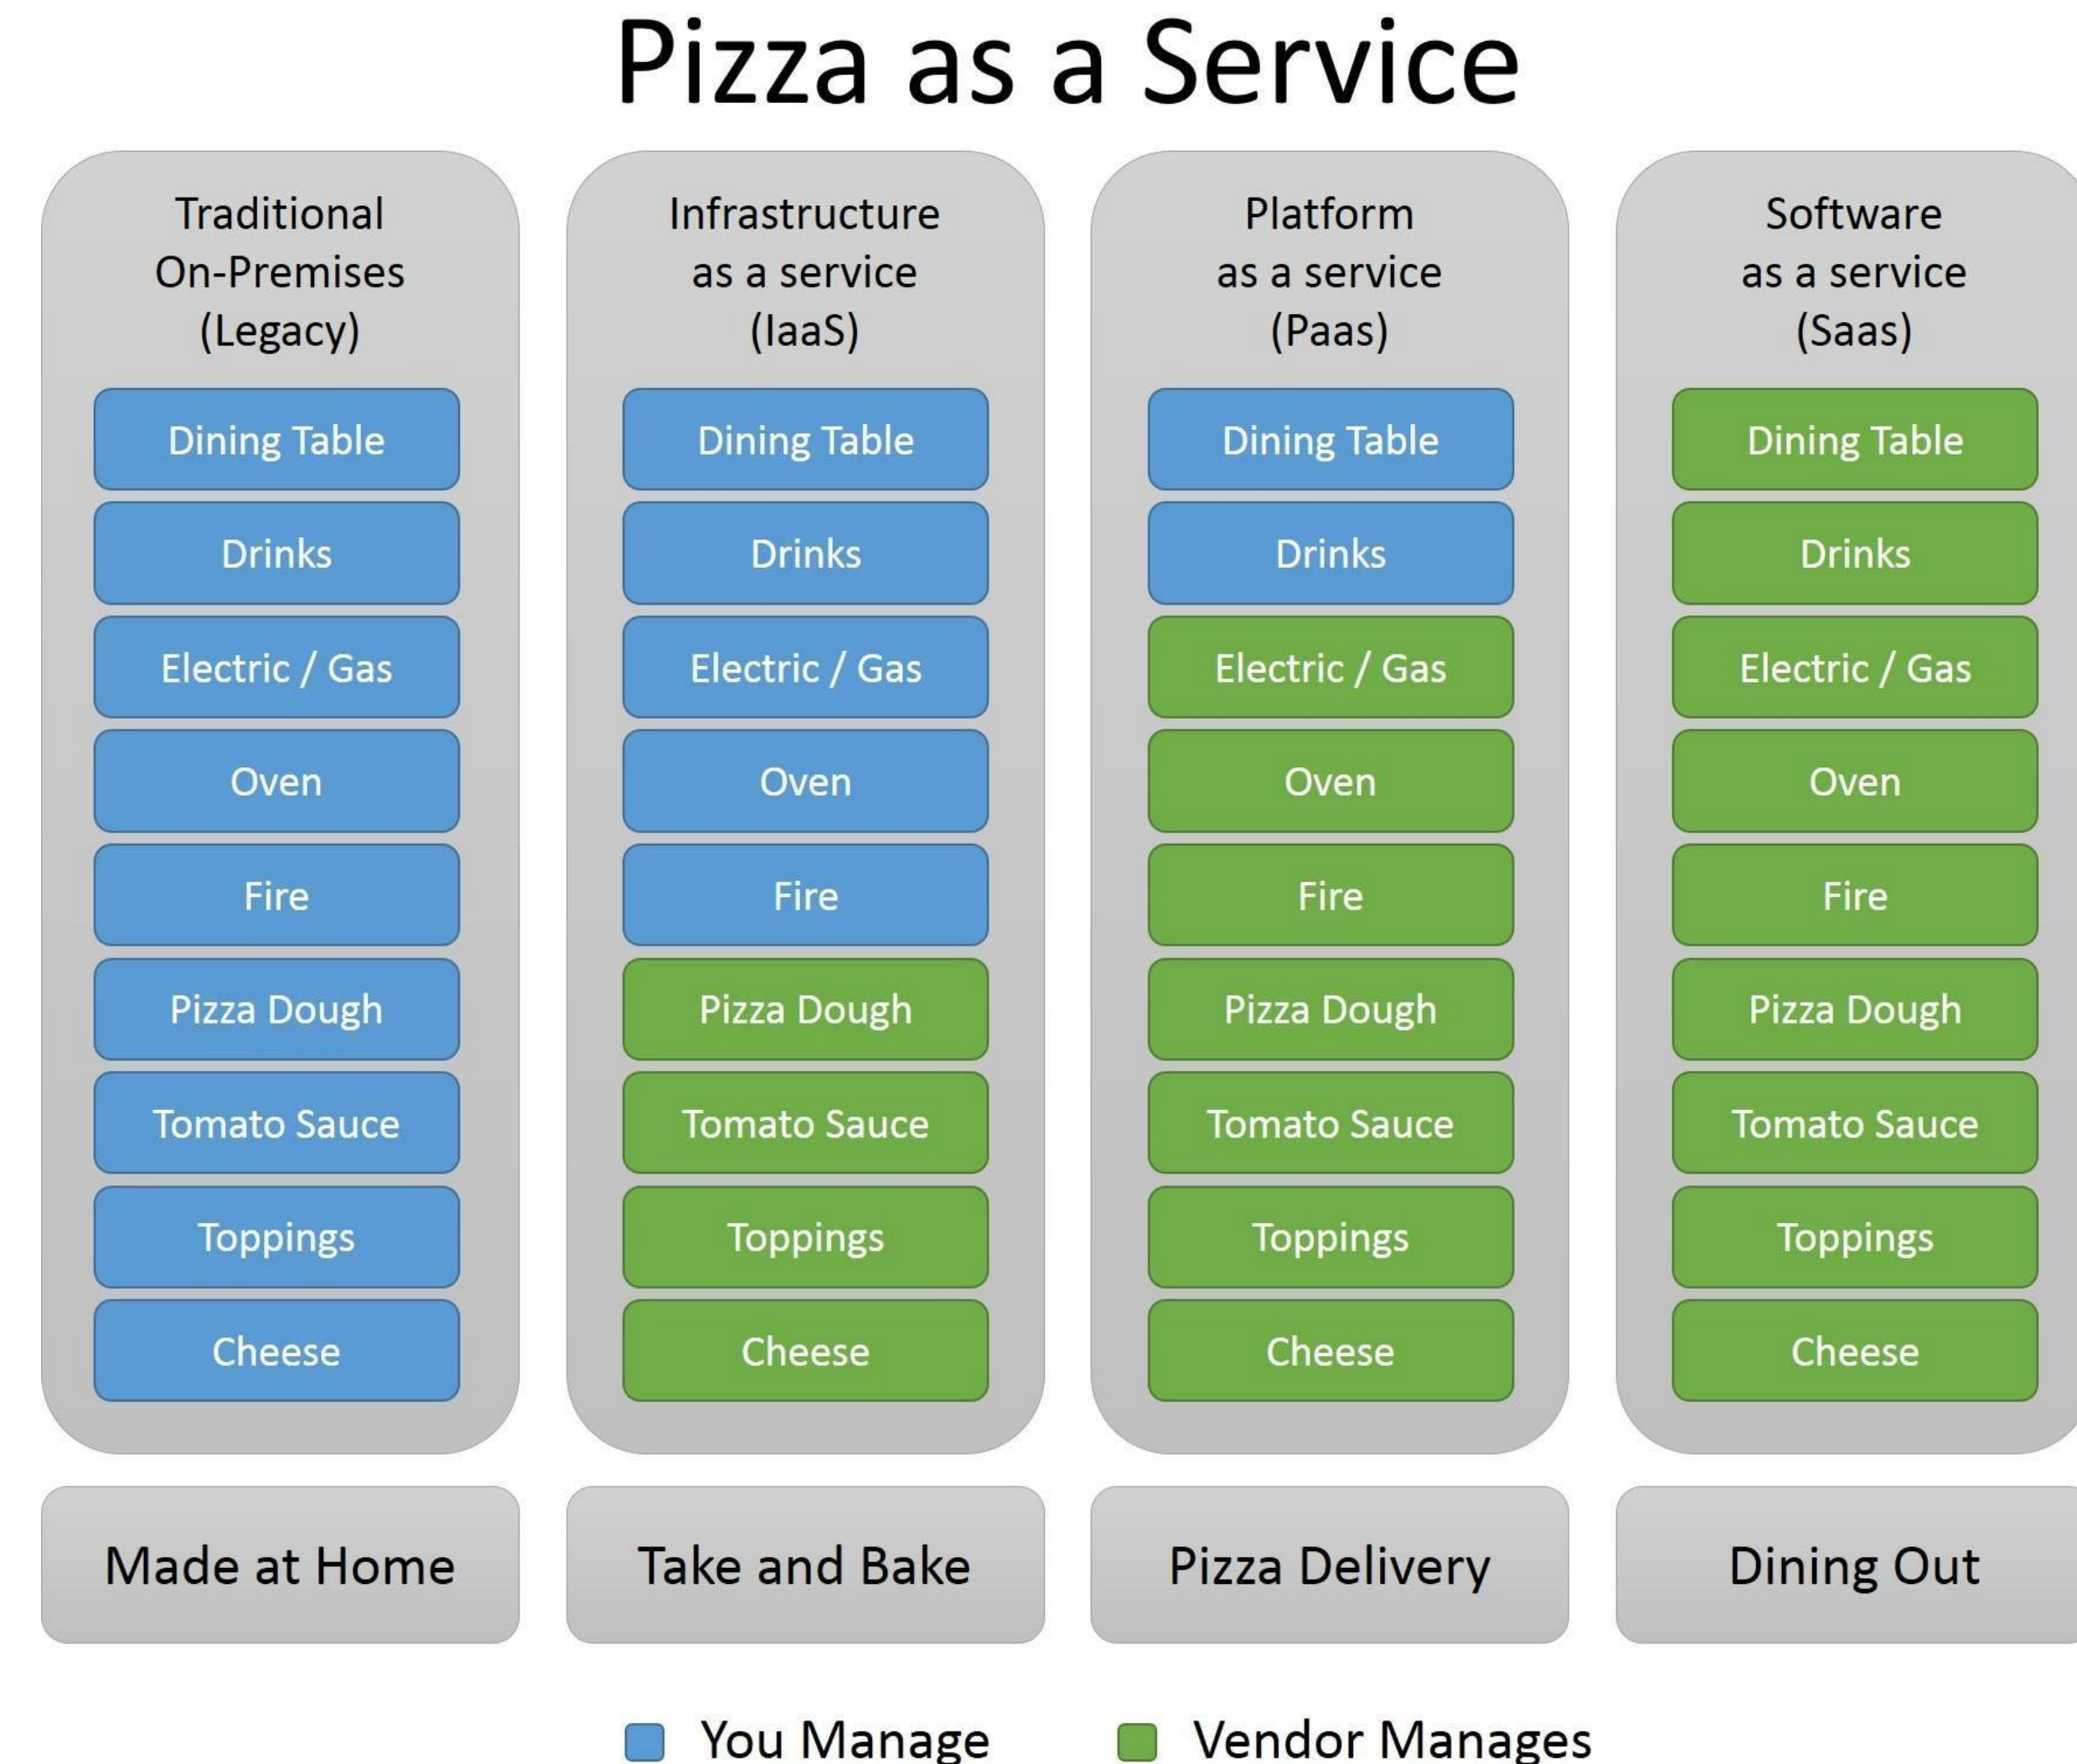
App Servers

Misc Services

Database servers

What parts of this infrastructure can be shared across different clients?

Shared infrastructure analogy: Pizza

- Four ways to get pizza: Make yourself, take and bake, delivery, dine out

- Vendor manages different levels of the stack, achieving economies of scale

- When would you choose one over the other?

Pizza as a Service

Traditional On-Premises (Legacy)	Infrastructure as a service (IaaS)	Platform as a service (Paas)	Software as a service (Saas)
Dining Table	Dining Table	Dining Table	Dining Table
Drinks	Drinks	Drinks	Drinks
Electric / Gas	Electric / Gas	Electric / Gas	Electric / Gas
Oven	Oven	Oven	Oven
Fire	Fire	Fire	Fire
Pizza Dough	Pizza Dough	Pizza Dough	Pizza Dough
Tomato Sauce	Tomato Sauce	Tomato Sauce	Tomato Sauce
Toppings	Toppings	Toppings	Toppings
Cheese	Cheese	Cheese	Cheese
Made at Home	Take and Bake	Pizza Delivery	Dining Out

■ You Manage   ■ Vendor Manages

Pizza as a Service — by Albert Barron (unlicensed?)

Shared infrastructure creates economies of scale

- At the physical level:

  - Multiple customers' physical machines in the same data center

  - Save on physical costs (centralize power, cooling, security, maintenance)

- At the physical server level:

  - Multiple customers' virtual machines in the same physical machine

  - Save on resource costs (utilize marginal computing capacity – CPUs, RAM, disk)

- At the application level:

  - Multiple customer's applications hosted in same virtual machine

  - Save on resource overhead (eliminate redundant infrastructure like OS)

- "Cloud" is the natural expansion of multi-tenancy at all levels

What is the infrastructure that can be shared?

- Our apps run on a "tall stack" of dependencies
- Old style: this full stack is self-managed
- Cloud providers offer products that manage parts of that stack for us:
  - "Infrastructure as a service"
  - "Platform as a service"
  - "Software as a Service"
  - Collectively called "X as a Service"

Traditional, on-premises computing
Application
Middleware
Operating System
Virtualization
Physical Server
Storage
Network
Physical data center

Self-managed

Platform-as-a-Service
Application
Middleware
Operating System
Virtualization
Physical Server
Storage
Network
Physical data center

Vendor-managed

Cloud services gives on-demand access to infrastructure, "as a service"

- Vendor provides a service catalog of "**X as a service**" abstractions that provide infrastructure as a service

- API allows us to provision resources on-demand

- Transfers responsibility for managing the underlying infrastructure to a vendor
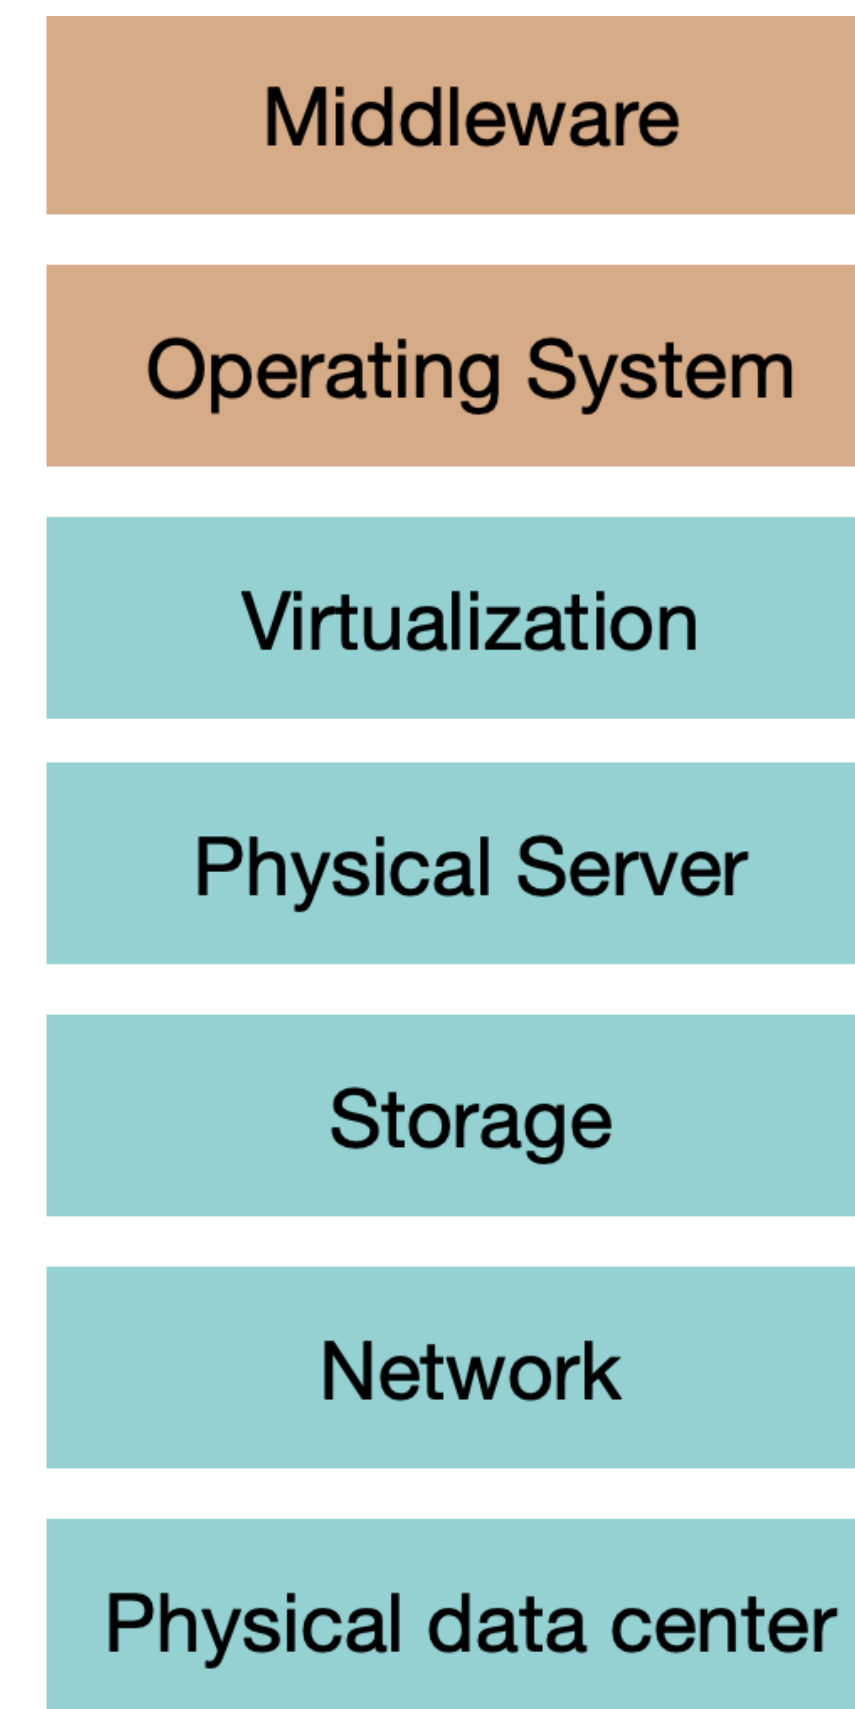
Cloud infrastructure scales elastically

- "Traditional" computing infrastructure requires capital investment
  - "Scaling up" means buying more hardware, or maintaining excess capacity for when scale is needed
  - "Scaling down" means selling hardware, or powering it off
- Cloud computing scales elastically:
  - "Scaling up" means allocating more shared resources
  - "Scaling down" means releasing resources into a pool
  - Billed on consumption (usually per-second, per-minute or per-hour)

Infrastructure as a Service: Virtual Machines

- Virtualize a single large server into many smaller machines
- Each VM runs its own OS
- OS limits resource usage and guarantees per-VM quality
- Administration responsibilities separated for physical machine vs virtual machines
- Examples:
-   Cloud: Amazon EC2, Google Compute Engine, Azure
-   On-Premises: VMWare, Proxmox

Middleware
Operating System
Virtualization
Physical Server
Storage
Network
Physical data center

IaaS

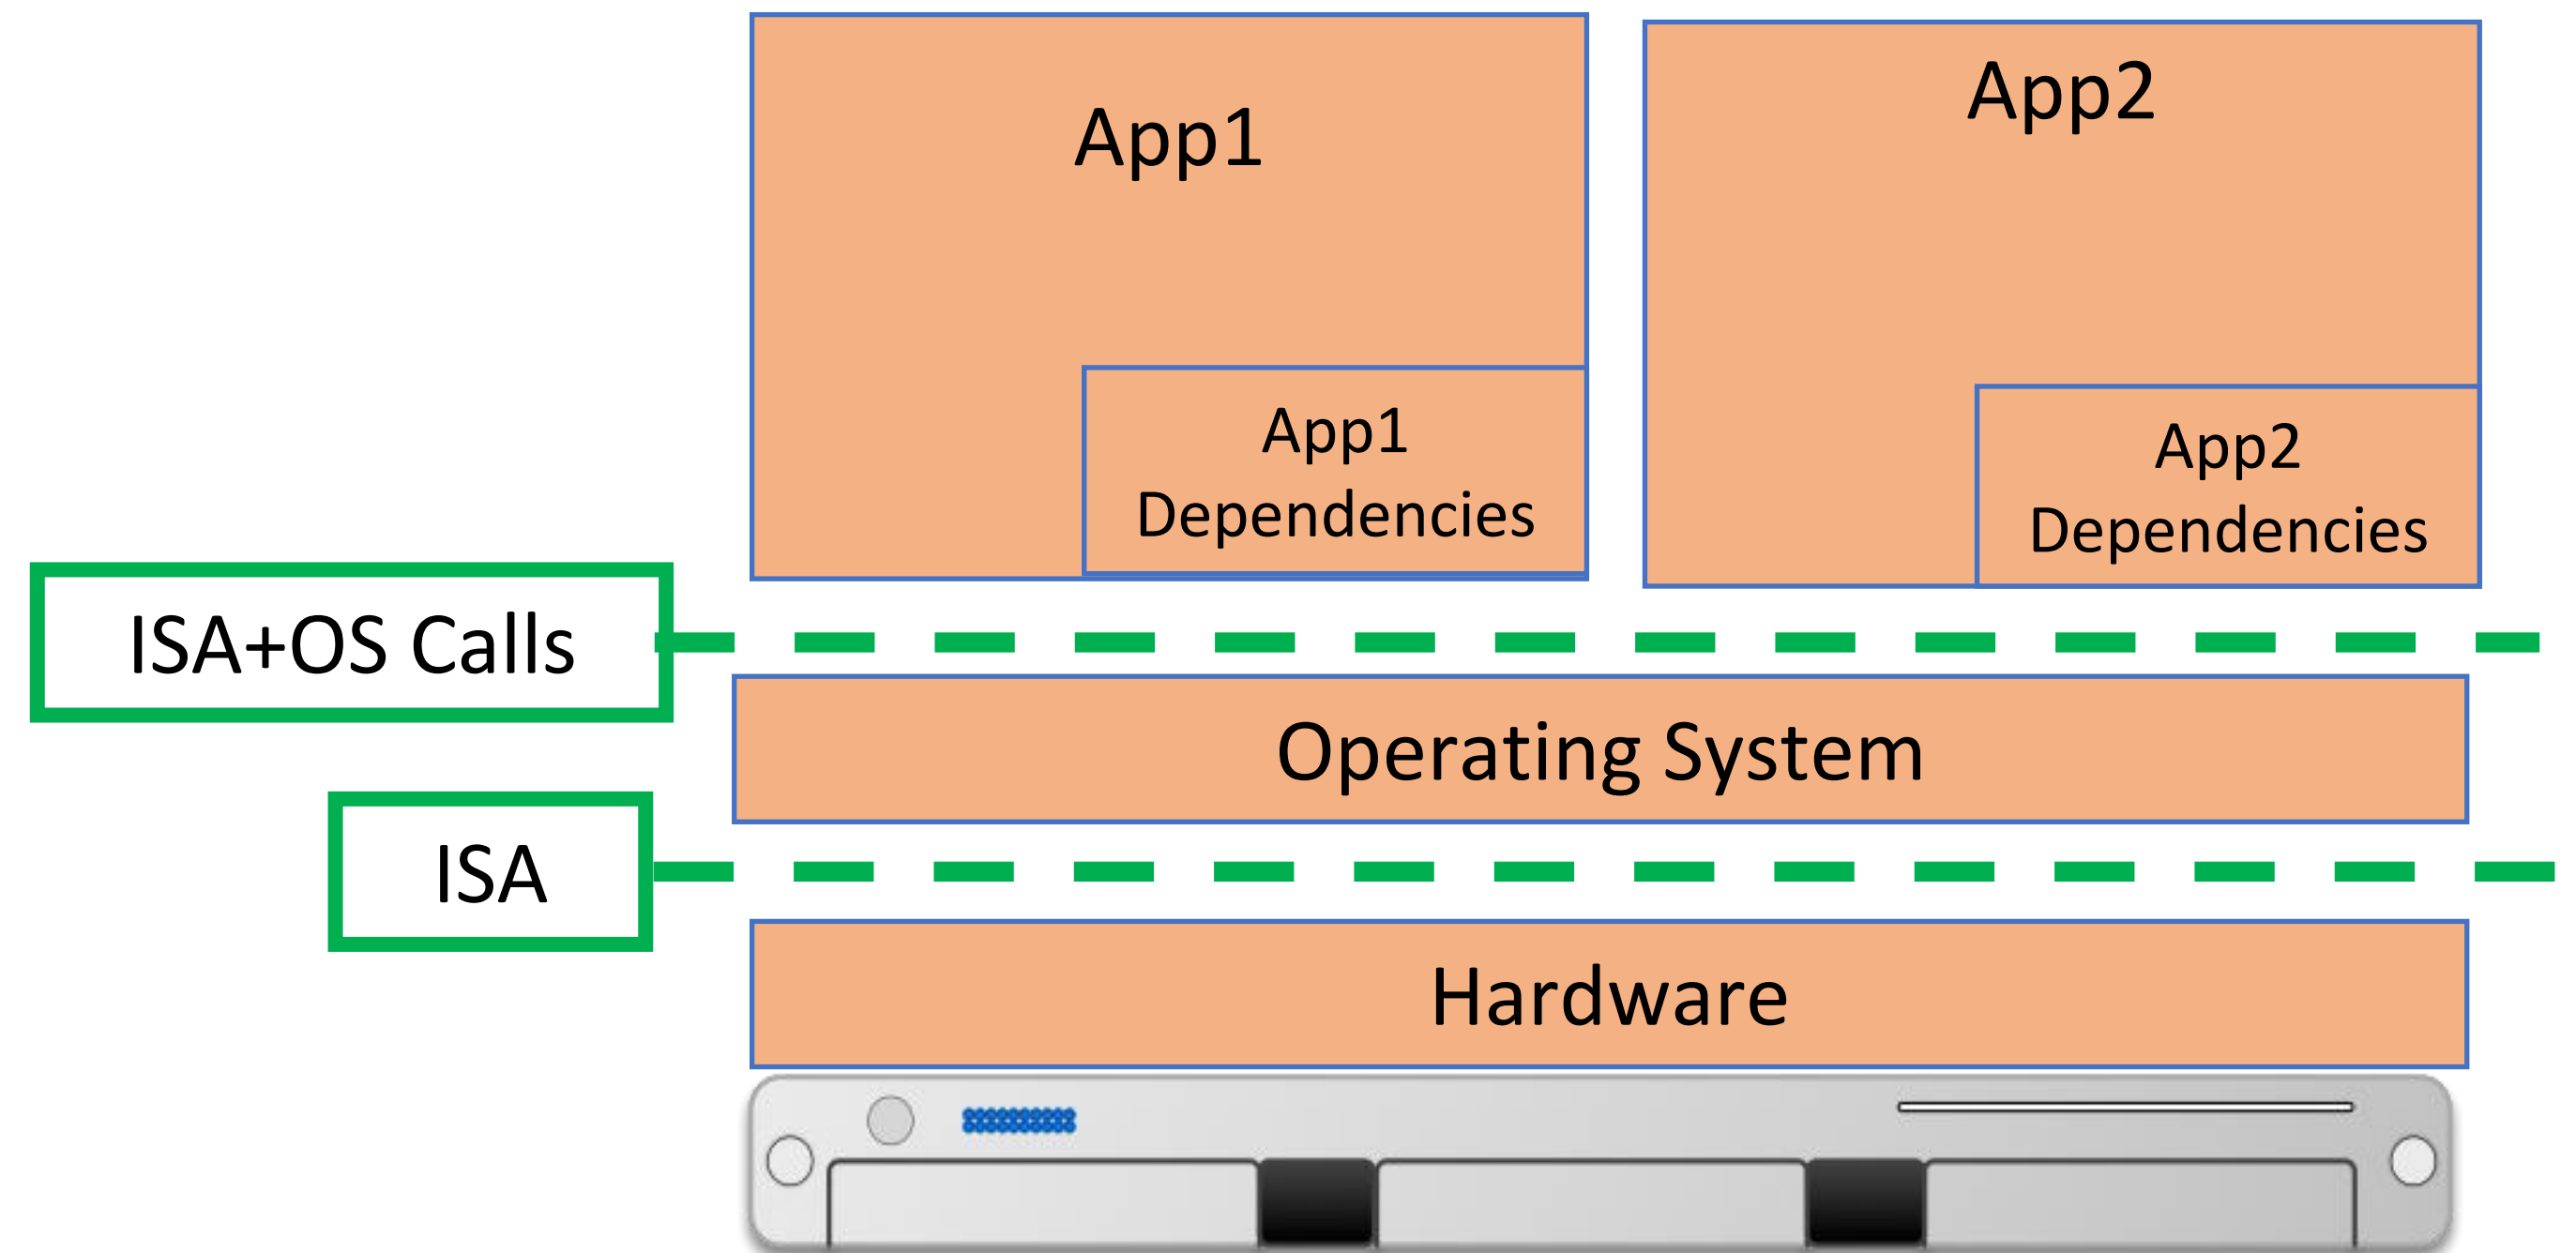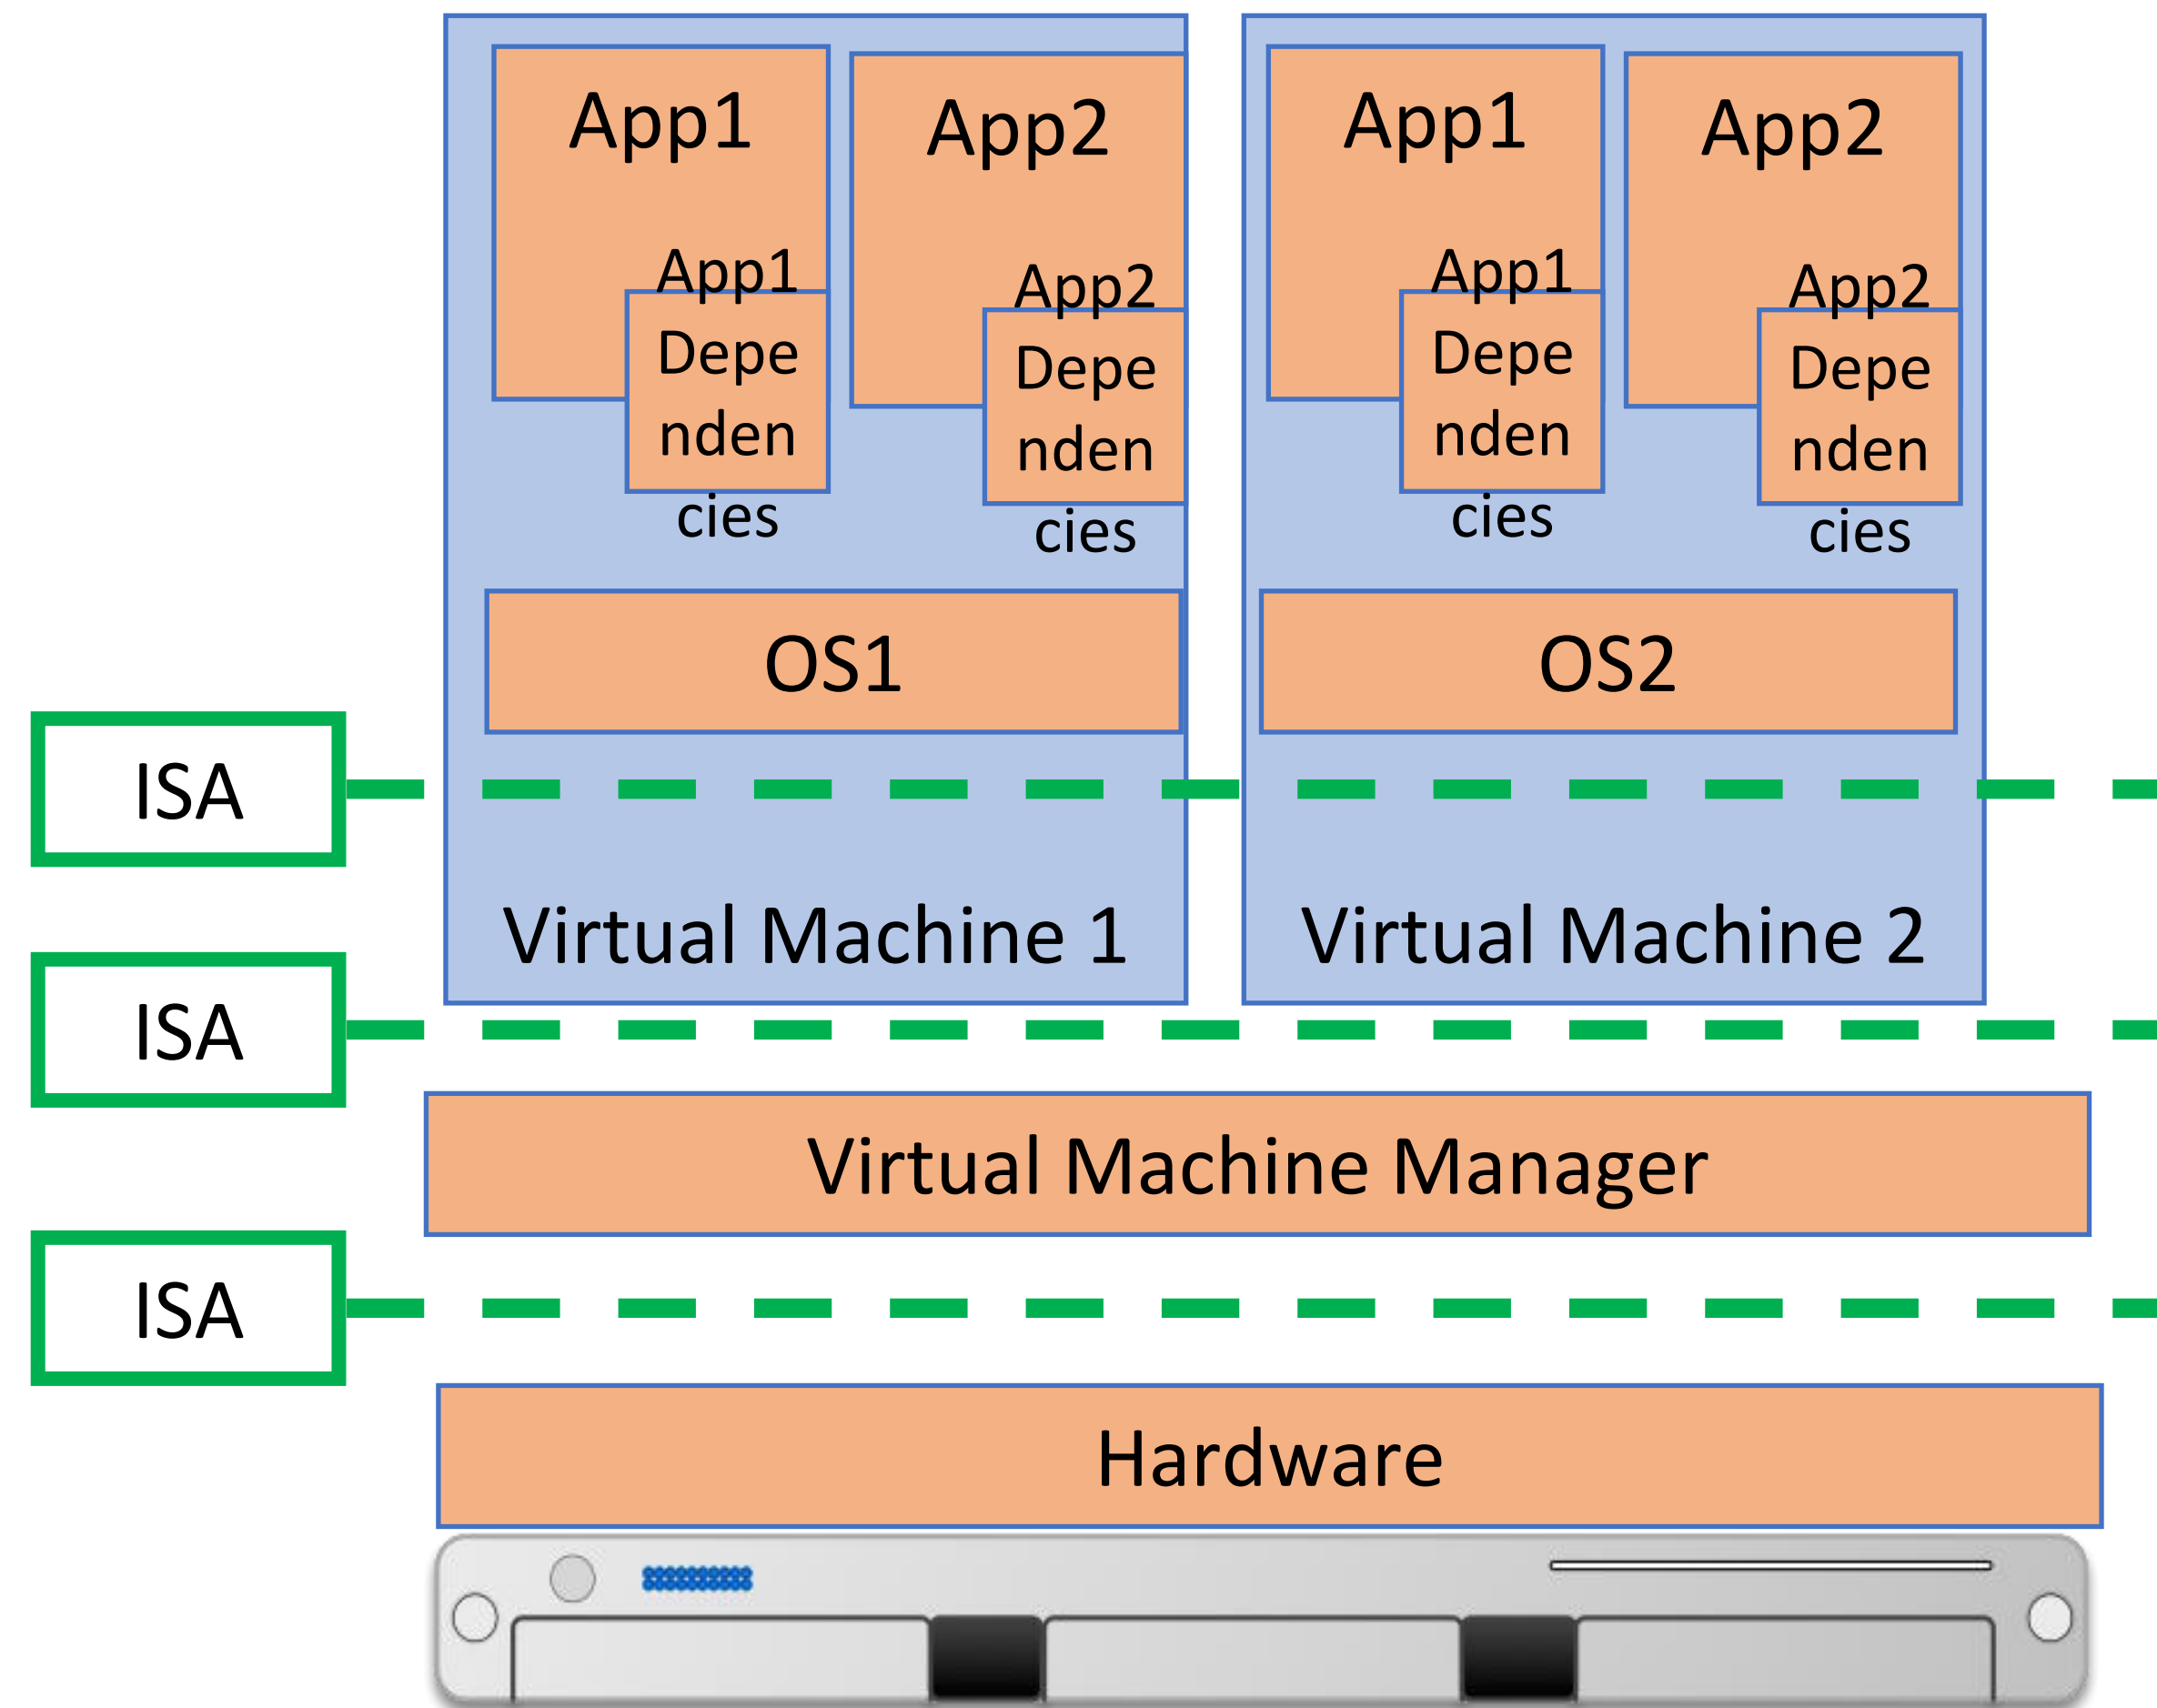| VM4	VM5	VM6
VM3		
VM1		VM2

Abstracted physical machine

The operating system allows several apps to share the underlying hardware

- The "instruction set" is an abstraction of the underlying hardware

- The operating system presents the same abstraction + OS calls.

A virtual machine layer allows several different operating systems to share the same hardware

Virtual Machines facilitate multi-tenancy

- Multi-Tenancy
  - Multiple customers sharing same physical machine, oblivious to each other

- Decouples application from hardware
  - virtualization service can provide "live migration" transparent to the operating system, maximizing utilization

- Faster to provision and release
  - VM v. physical machines == ~mins v. ~hours (days?)

Virtual Machines to Containers

- Each VM contains a full operating system

- What if each application could run in the same (overall) operating system? Why have multiple copies?

- Advantages to smaller apps:

  - Faster to copy (and hence provision)

  - Consume less storage (base OS images are usually 3-10GB)

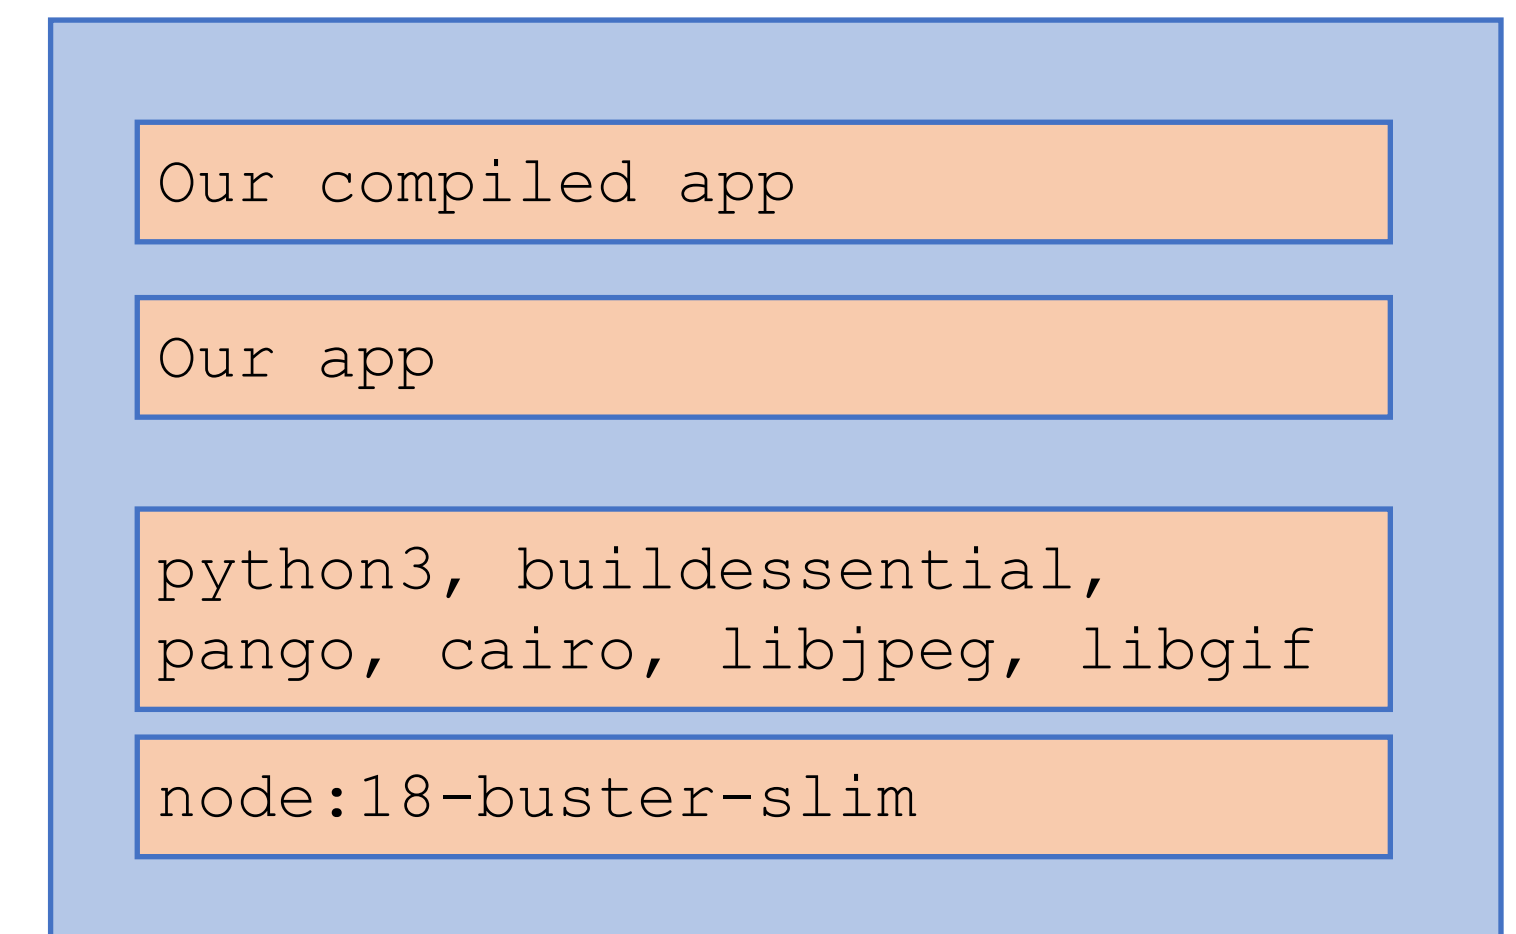# Containers run layered images, reducing storage space

- Images are defined programmatically as a series of "build steps" (e.g. Dockerfile)

- Each step in the build becomes a "layer"

- Built layers can be shared and cached

- To run a container, the layers are linked together with an "overlay" filesystem

```
FROM node:18-buster-slim
RUN apt-get update && apt-get install python3
build-essential libpango1.0-dev libcairo2-dev
libjpeg-dev libgif-dev -y

RUN mkdir -p /usr/src/app
WORKDIR /usr/src/app
COPY ./ /usr/src/app

RUN npm ci
RUN npm run build
CMD [ "npm", "start" ]
```

Example image specification (Dockerfile)

| Our compiled app |
| Our app |
| python3, buildessential, pango, cairo, libjpeg, libgif |
| node:18-buster-slim |

Example image, with layers shown

Containers run layered images, reducing storage space

- Many images may share the *same* lower layers (e.g. OS, NodeJS, some system dependencies)

- Layers are shared between images

- Multi-tenancy: *N* running containers only require *one* copy of each layer (they are read-only)

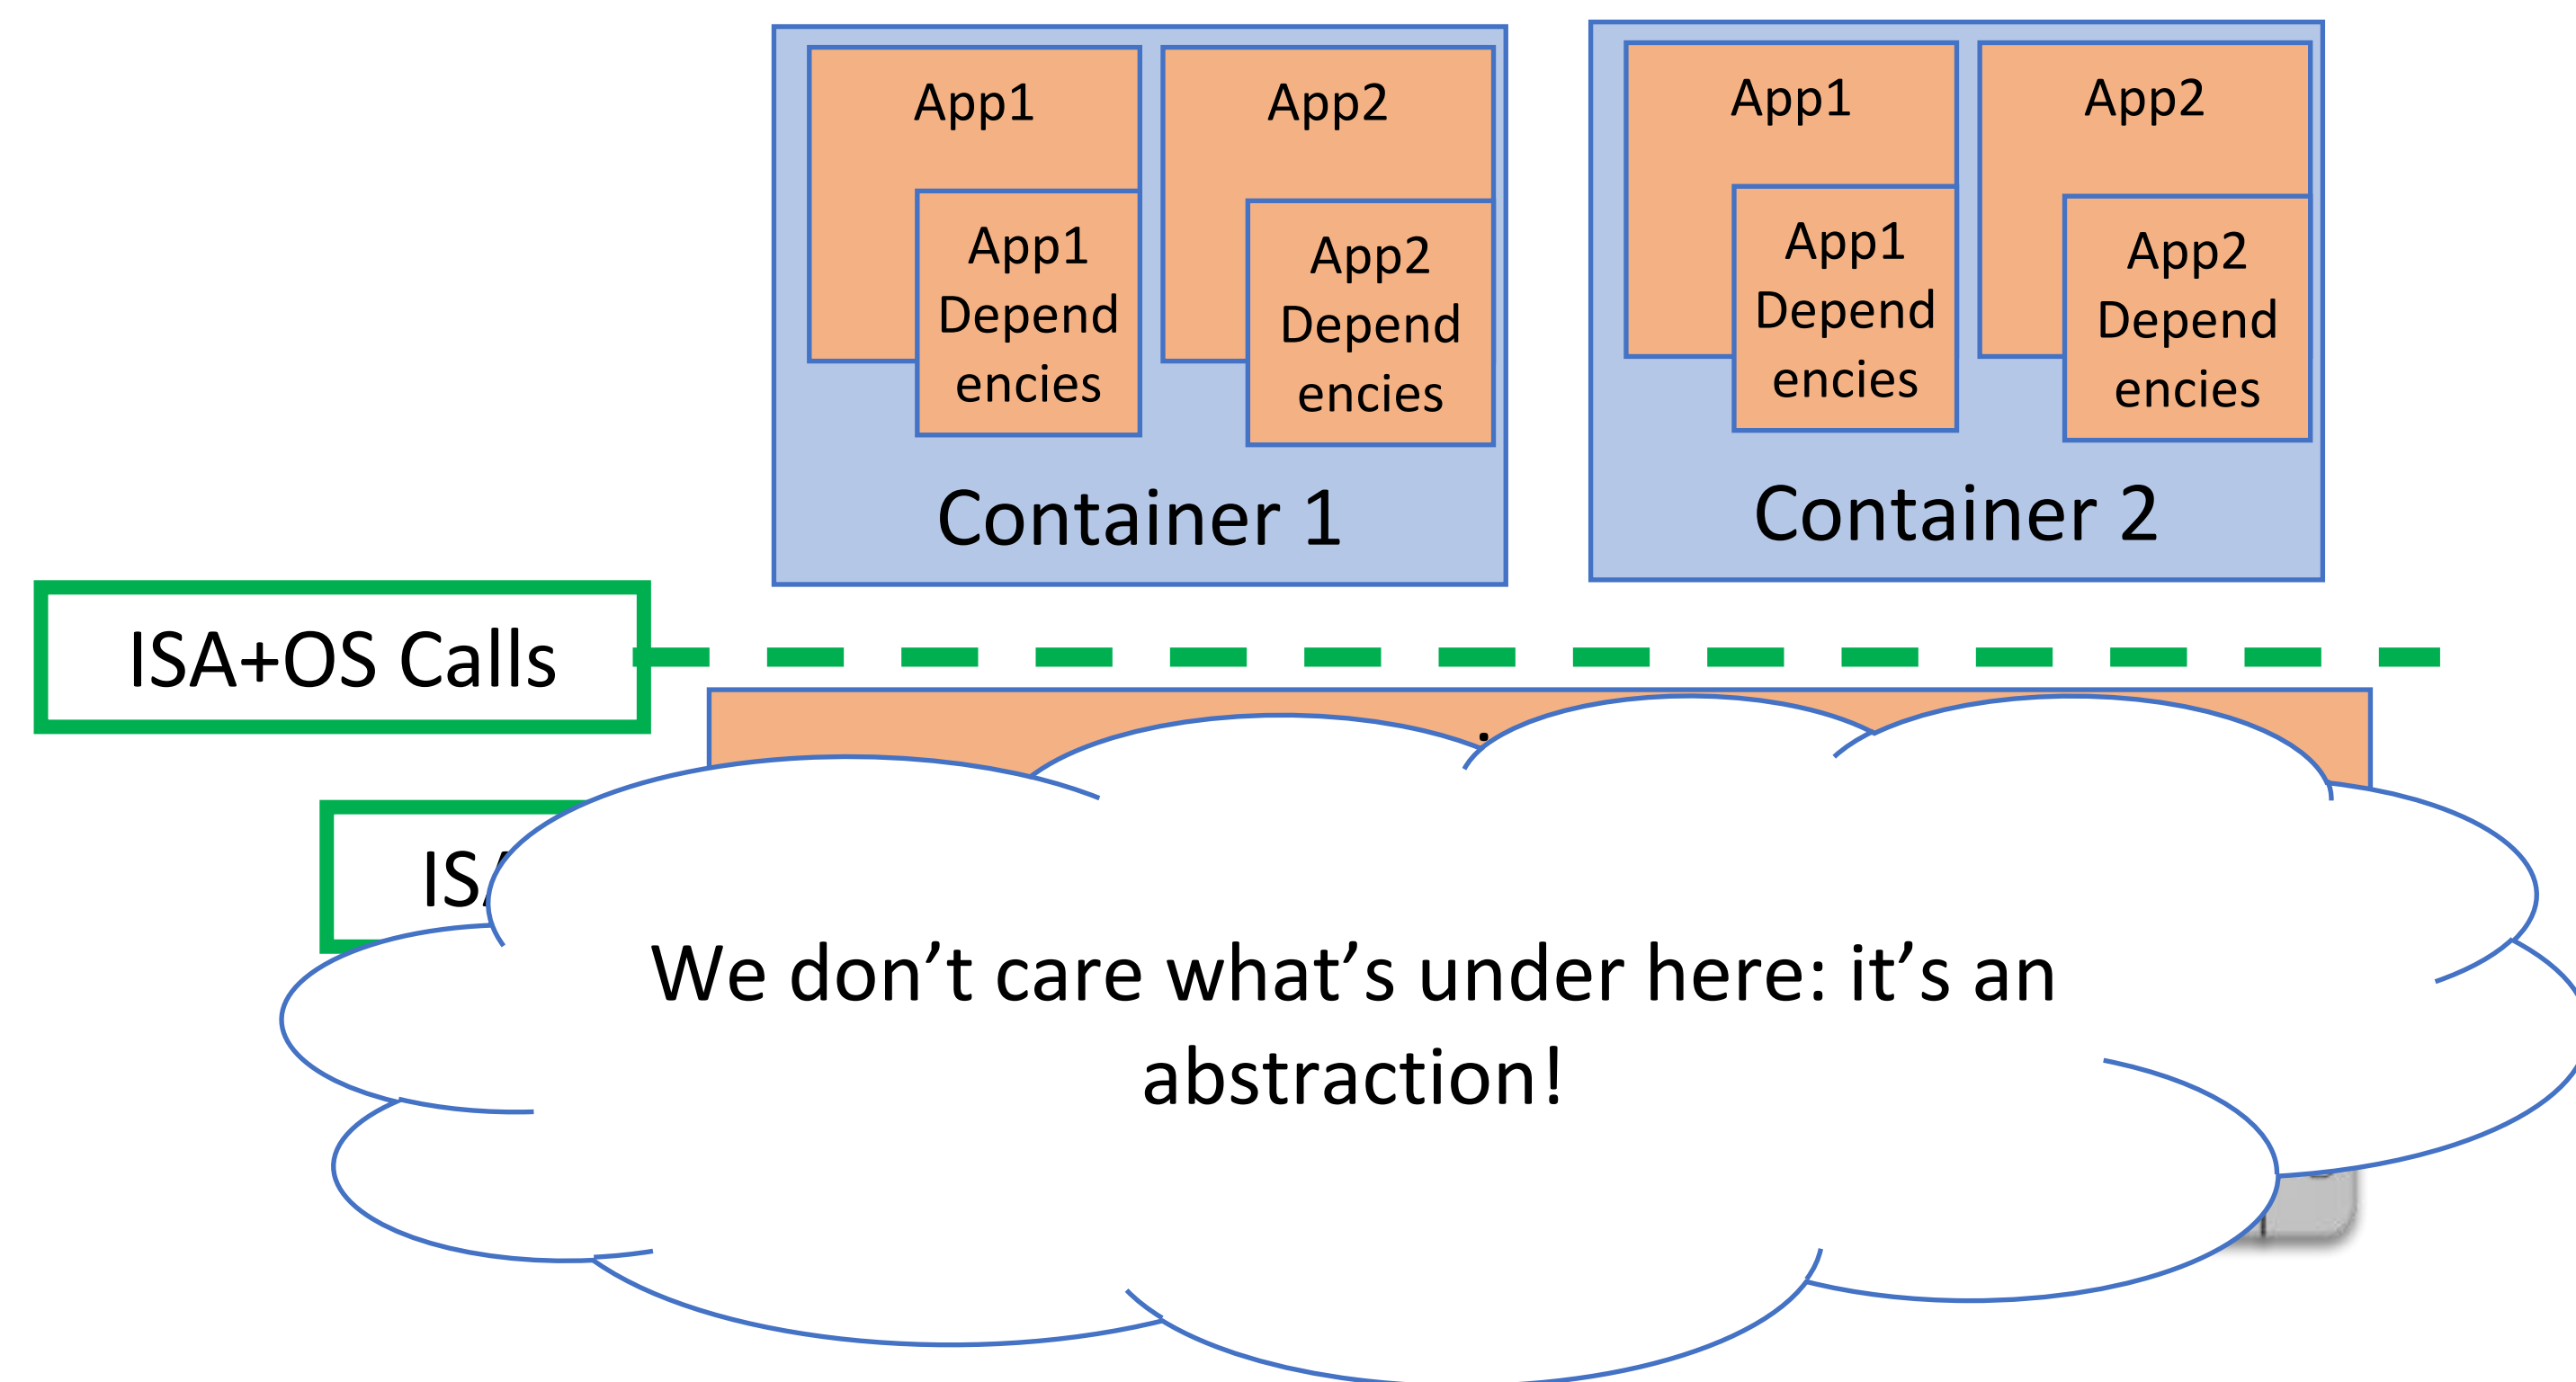| Orion's compiled app	Ripley's compiled app
Orion's app	Ripley's app
python3, buildessential, pango, cairo, libjpeg, libgif	python3, buildessential, pango, cairo, libjpeg, libgif
node:18-buster-slim	node:18-buster-slim

Two images, sharing two layers

A container contains your apps and all their dependencies

- Each application is encapsulated in a "lightweight container," includes:
  - System libraries (e.g. glibc)
  - External dependencies (e.g. nodejs)
- "Lightweight" in that container images are smaller than VM images - multi tenant containers run in the OS
- Cloud providers offer "containers as a service" (Amazon AWS Fargate, Azure Kubernetes, Google Kubernetes)

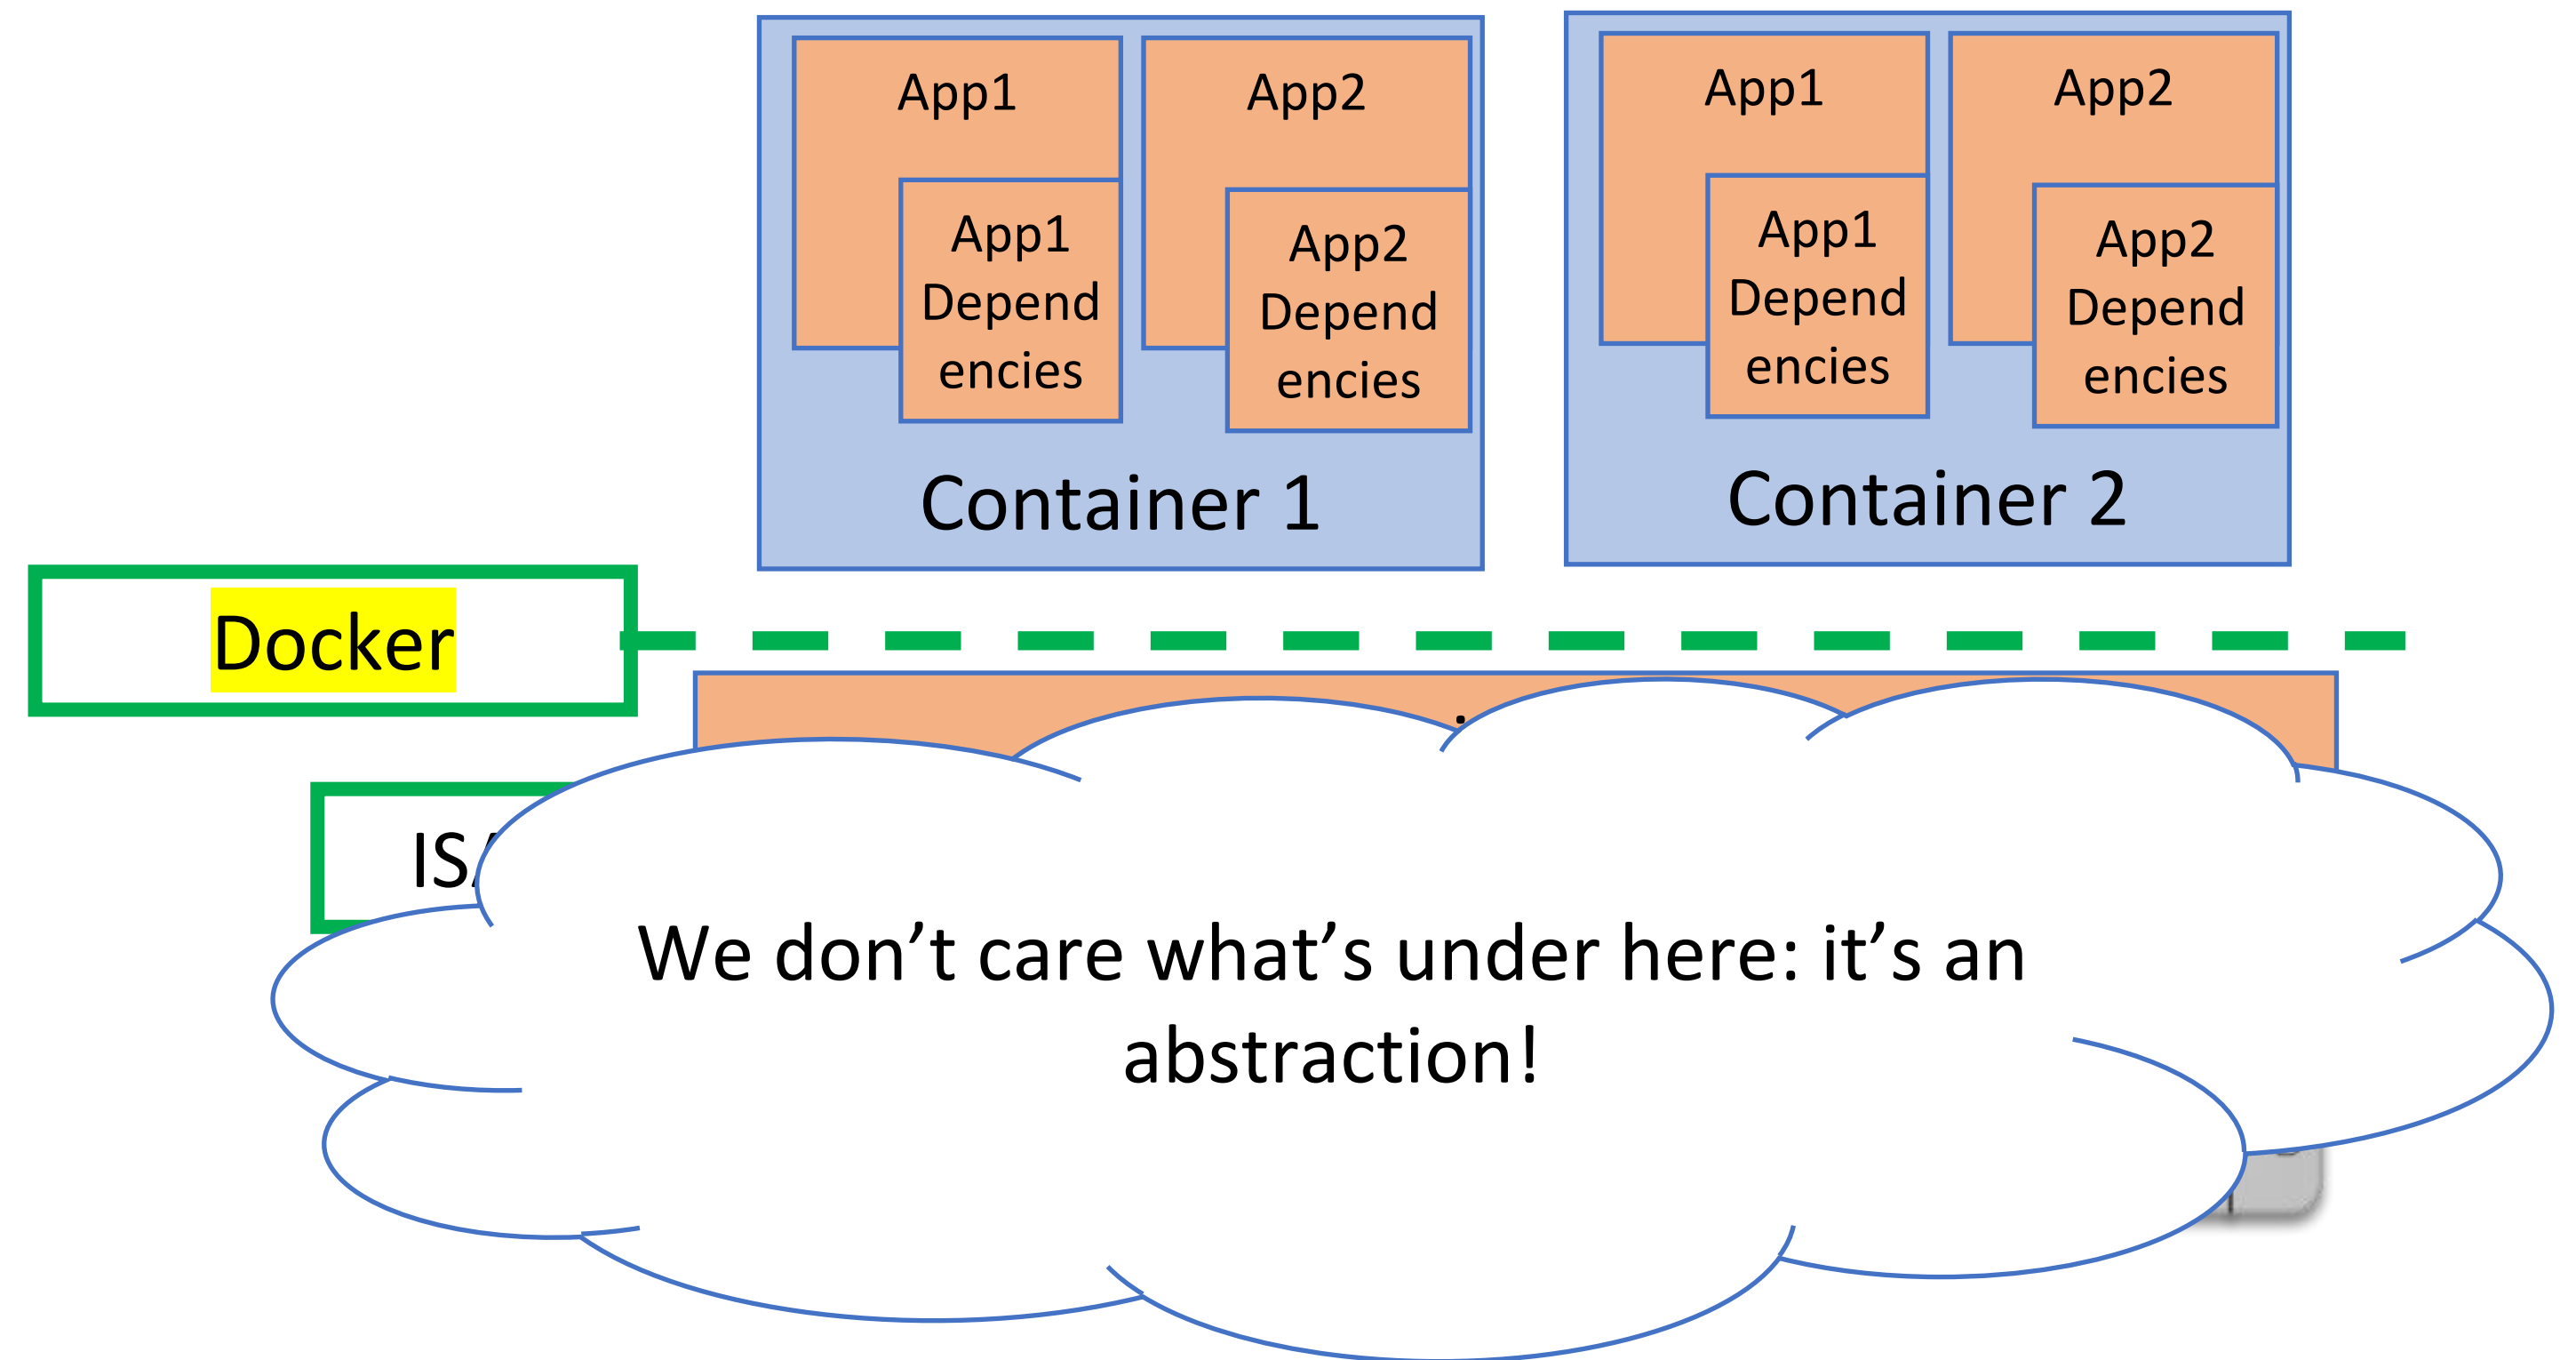# XaaS: Containers as a Service

- Vendor supplies an on-demand instance of an operating system

  - e.g.: Linux version NN

- Vendor is free to implement that instance in a way that optimizes costs across many clients.

Docker is the prevailing container platform

- Docker provides a standardized interface for your container to use

- Many vendors will host your Docker container

- An open standard for containers also exists ("OCI")

App1

App2

App1 Depend encies

App2 Depend encies

Container 1

App1

App2

App1 Depend encies

App2 Depend encies

Container 2

Docker

ISA

We don't care what's under here: it's an abstraction!

Tradeoffs between VMs and Containers

- Performance is comparable
- Each VM has a copy of the OS and libraries
  - Higher resource overhead
  - Slower to provision
  - Support for wider variety of OS's
- Containers are "lightweight"
  - Lower resource overhead
  - Faster to provision
  - Potential for compatibility issues, especially with older software

# Platform-as-a-Service: vendor supplies OS + middleware

- Middleware is the stuff between our app and a user's requests:
  - Content delivery networks: Cache static content
  - Web Servers: route client requests to one of our app containers
  - Application server: run our handler functions in response to requests from load balancer
  - Monitoring/telemetry: log requests, response times and errors
- Cloud vendors provide managed middleware platforms too: "Platform as a Service"

Internet

Clients

Content Delivery Network

Web Servers

App Servers

Monitoring/Telemetry

Database servers

PaaS is often the simplest choice for app deployment

- **Platform-as-a-Service** provides components most apps need, fully managed by the vendor: load balancer, monitoring, application server

- Some PaaS run your app in a container: Heroku, AWS Elastic Beanstalk, Google App Engine, Railway, Vercel...

- Other PaaS run your apps as individual functions/event handlers: AWS Lambda, Google Cloud Functions, Azure Functions

- Other PaaSs provide databases and authentication, and run your functions/event handlers: Google Firebase, Back4App
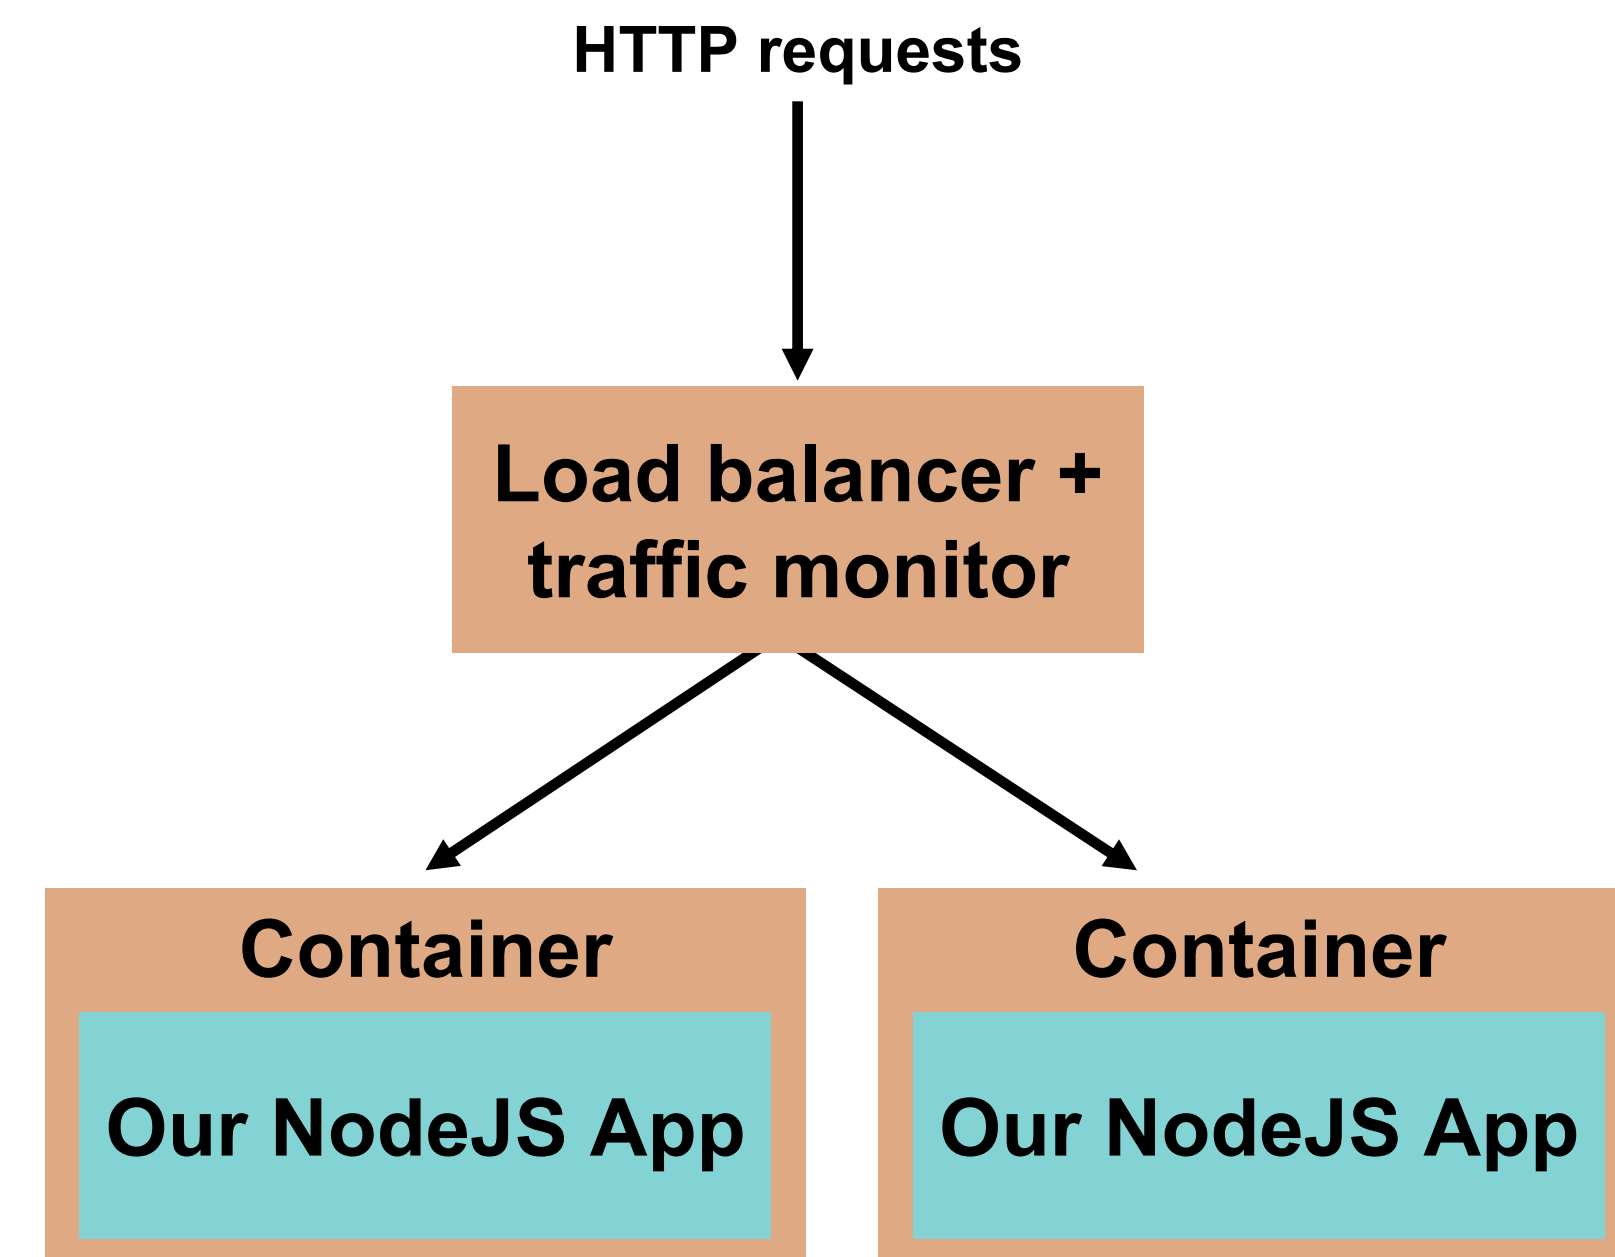
Application
Middleware
Operating System
Virtualization
Physical Server
Storage
Network
Physical data center

PaaS

PaaS in the style of Heroku runs containers

- Takes a web app as input
  - Provide an entry point to code, e.g. "npm start", or optionally, a container specification
- Hosts web app at chosen URL, can scale resources up/down on-demand
  - Load balancer fully managed by Heroku, scaling transparent
  - Auto-scale down to use no resources, spins up container on reception of a request
  - Dashboard for monitoring/reporting
- Newcomers provide similar functionality (Vercel, Railway, etc)
- Host PaaS on-premises, too (Caprover)

HTTP requests

Load balancer + traffic monitor

Container

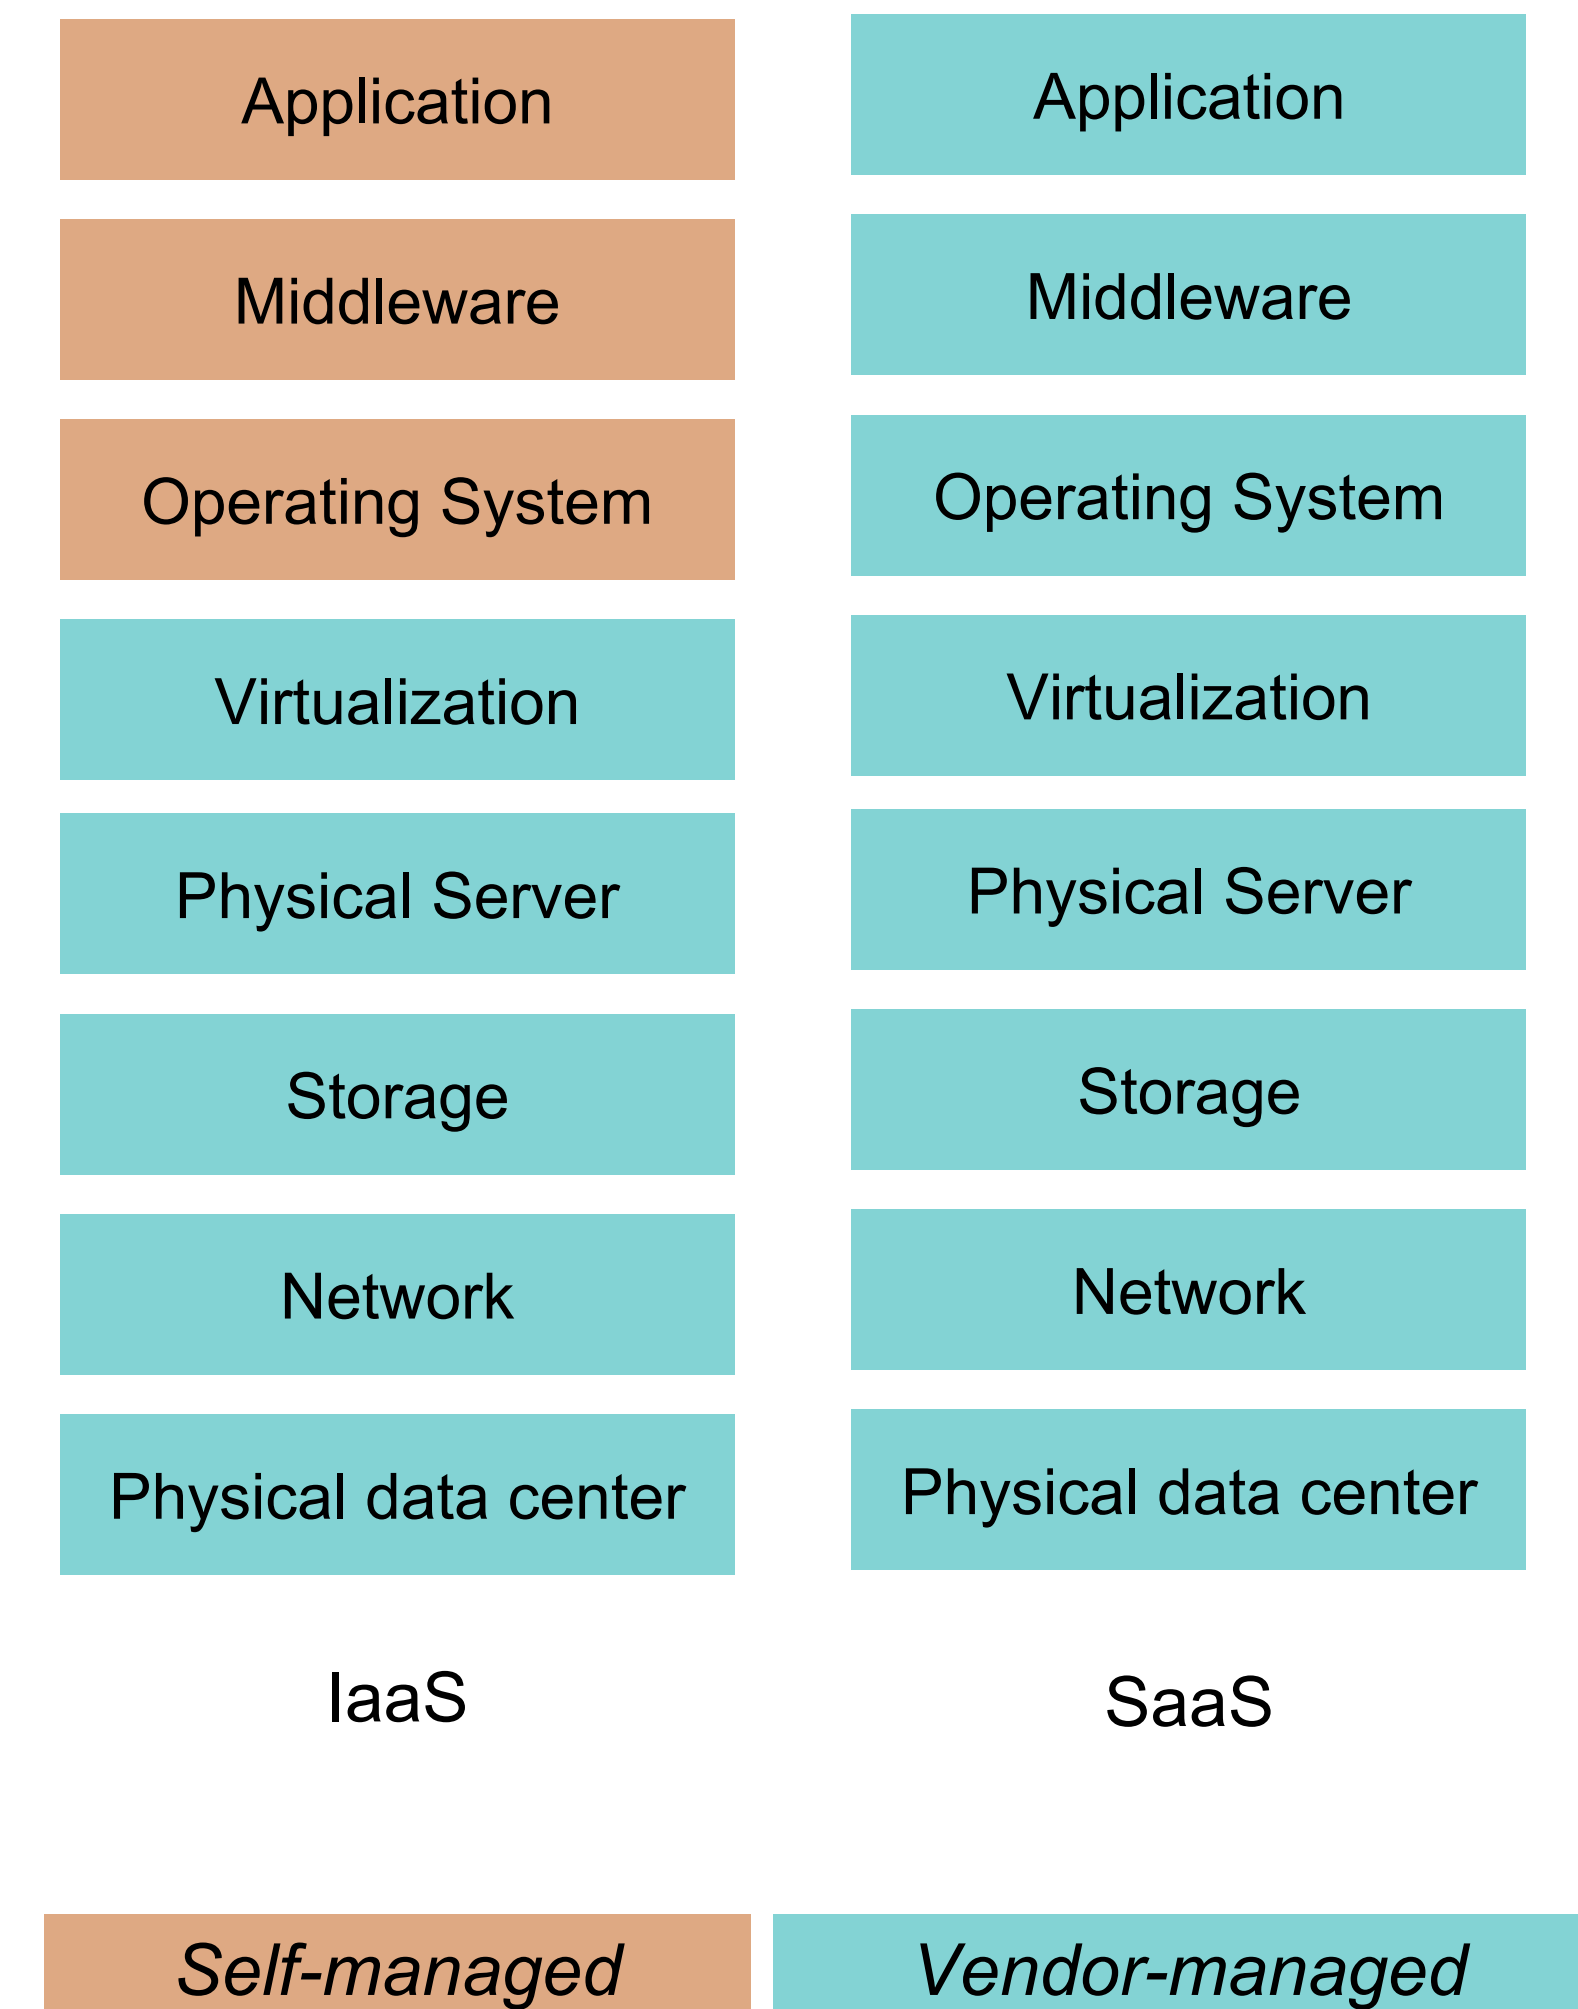**Our NodeJS App**

Container

Our NodeJS App

Render.com uses a mixed strategy for implementation

- Web services are run in containers on their app servers.

  - Load balancers, dashboards, etc. are shared by all users

  - Other services: eg Redis ("Software as a Service")

- "Static Web Sites" are hosted as static files on their CDN.

  - minimizes load times for users world-wide.

Software as a Service adds more vendor-managed apps

- Providers may also develop custom software offered only as a service

- Examples:

  - PostgreSQL (open source)

  - Twilio Programmable Video (proprietary chat)

IaaS	SaaS
Application	Application
Middleware	Middleware
Operating System	Operating System
Virtualization	Virtualization
Physical Server	Physical Server
Storage	Storage
Network	Network
Physical data center	Physical data center

Self-managed   *Vendor-managed*

Self-managed vs Vendor-managed Infrastructure Tradeoffs

- Consider who manages each tier in the stack
- Benefits to vendor-managed options:
  - More ways to reduce resource consumption, improve resource utilization
  - Less management burden
  - Less capital investment, more flexibility in scaling
- Benefits to self-managed options:
  - Greater flexibility to migrate between software platforms
  - More capital investment, potentially less operating expense

Traditional, on-premises computing	SaaS
Application	Application
Middleware	Middleware
Operating System	Operating System
Virtualization	Virtualization
Physical Server	Physical Server
Storage	Storage
Network	Network
Physical data center	Physical data center

Self-managed	*Vendor-managed*

Cloud Infrastructure can be better for variable workloads

- Consider:
  - Does your workload benefit from ability to scale up or down?
  - Variable workloads have different demands over time (most common)
  - Constant workloads require sustained resources (less common)
- Example:
  - Need to run 300 VMs, each 4 vCPUs, 16GB RAM
- Private cloud:
  - Dell PowerEdge Pricing (AMD EPYC 64 core CPUs)
  - 7 servers, each 128 cores, 512GB RAM, 3 TB storage = $162,104
- Public cloud:
  - Amazon EC2 Pricing (M7a.xlarge instances, $0.153/VM-hour)
  - 10 VMs for 1 year + 290 VMs for 1 month: $45,792.90
  - 300 VMs for 1 year: $402,084.00

Public clouds are not the only option

- "Public" clouds are connected to the internet and available for anyone to use

  - Examples: Amazon, Azure, Google Cloud, DigitalOcean

- "Private" clouds use cloud technologies with on-premises, self-managed hardware

  - Cost-effective when a large scale of baseline resources are needed

  - Example management software: OpenStack, VMWare, Proxmox, Kubernetes

- "Hybrid" clouds integrate private and public (or multiple public) clouds

  - Effective approach to "burst" capacity from private cloud to public cloud

"X as a Service" offers several abstractions to choose from depending on your needs

- Vendor manages different levels of the stack, achieving economies of scale

- When would you choose one over the other?

- Explore some options at https://comparecloud.in/

Traditional, on-premises computing	IaaS	PaaS	SaaS
Application	Application	Application	Application
Middleware	Middleware	Middleware	Middleware
Operating System	Operating System	Operating System	Operating System
Virtualization	Virtualization	Virtualization	Virtualization
Physical Server	Physical Server	Physical Server	Physical Server
Storage	Storage	Storage	Storage
Network	Network	Network	Network
Physical data center	Physical data center	Physical data center	Physical data center

Self-managed

Vendor-managed

Review

- You should now be able to…

  - Explain what "cloud" computing is and why it is important

  - Explain why shared infrastructure is important in cloud computing

  - Describe the difference between virtual machines and containers

  - Discuss trade-offs that you might consider for self or vendor-managed platforms